

**SOLANGE: A Provenance-Complete, Enterprise-Architected QC·AI·HPC Orchestration
Platform for Non-Druggable NSCLC Tumor-Suppressor Mutations**

A dissertation submitted to the faculty of Claremont Graduate University
in partial fulfillment of the requirements for the degree of Doctor of Philosophy

by

Doron Cohen

Claremont Graduate University

2026

Approval of the Dissertation Committee

This dissertation has been duly read, reviewed, and critiqued by the Committee listed below, which hereby approves the manuscript of Doron Cohen as fulfilling the scope and quality requirements for meriting the degree of Doctor of Philosophy.

Prof. Itamar Shabtai, Chair

Claremont Graduate University

[Committee Member]

[Committee Member]

Abstract of the Dissertation

SOLANGE: A Provenance-Complete, Enterprise-Architected QC·AI·HPC Orchestration Platform for Non-Druggable NSCLC Tumor-Suppressor Mutations

by

Doron Cohen

Claremont Graduate University: 2026

Approximately 60% of non-small cell lung cancer (NSCLC) patients harbor mutations for which no approved targeted therapy exists. Among the most clinically significant are structural TP53 DNA-interface mutants — particularly p.Cys275Phe (C275F) — which disrupt the DNA-binding interface in ways that are intractable to classical computational drug discovery methods at the accuracy required for reliable candidate ranking. The fermionic Hamiltonian eigenvalue problem at the C275F active-space boundary belongs to a family of problems for which quantum simulation algorithms (Abrams & Lloyd, 1999; Aspuru-Guzik et al., 2005) provide efficient solutions given an adequate initial state, placing it within BQP under the standard complexity-theoretic framing of quantum chemistry. This is an operational BQP-relevance claim, not a formal separation: ground-state finding is QMA-hard in the general case, and BQP containment is conditional on sufficient initial-state overlap (§1.2). Under that framing the dual operational condition holds that (i) the best known classical methods do not achieve the required accuracy at this scale within tractable resources, and (ii) the problem admits a known BQP-class algorithmic solution in the selected operational regime.

This dissertation proposes, designs, implements, and evaluates an orchestration platform — hardware-agnostic at the workflow level; vendor-specific at the compilation and error-mitigation layer — that integrates quantum computing, high-performance classical computing, and artificial intelligence into a unified, reproducible, and FDA 21 CFR Part 11-compliant workflow for simulating non-druggable NSCLC tumor suppressor mutation targets. The primary

research artifact is a working prototype platform grounded in real patient NGS data, validated through proof-of-concept execution on a quantum simulator backend.

Three original IT contributions are claimed: a BQP-relevant classification framework for non-druggable NSCLC targets, a Part-11-governed QC·AI·HPC orchestration architecture for oncology mutation simulation, and the first compliance framework (as of June 2026) applying FDA 21 CFR Part 11 to probabilistic quantum-computed molecular binding data — including ML-decoder provenance (P9) for emerging Ising-class learned QEC. The contribution of this dissertation is the architecture — the system that makes quantum simulation of non-druggable tumor suppressor mutations accessible, reproducible, and governed. Full-scale execution on IBM Heron r3 quantum hardware is identified as a subsequent research phase, enabled by the architecture designed here.

Current targeted therapies, even where they exist, confront an additional structural barrier: de novo resistance. For the majority of NSCLC patients in the non-druggable class represented by Table 2, no first-line targeted agent has ever been identified — and therefore no resistance trajectory has even begun. HIV provides the closest historical analogue at scale: a uniformly fatal disease with no therapeutic entry point, transformed by computational structural biology and iterative molecular simulation into a manageable chronic condition. The transformation did not happen because a single drug was found — it happened because the information systems infrastructure for drug-discovery iteration was built: governed, reproducible, and grounded in molecular simulation. This dissertation proposes quantum simulation as the foundational computational step toward an analogous transformation for non-druggable NSCLC tumor-suppressor mutations — not as a therapeutic claim, but as the information systems architecture that makes iterative computational drug discovery possible for the first time in this class.

Table of Contents

§ 01 Problem Statement — Three Gaps, One System.
§ 02 Research Questions — Four Questions.
§ 03 Literature Review & Research Gaps — What Exists, What's Missing.
§ 04 Theoretical Framework — The Intellectual Scaffolding.
§ 05 Research Methodology — Five Phases.
§ 06 Original Contributions — Three Claims.
§ 07 Preliminary Work — A Deployed Prototype.
§ 08 Limitations & Scope — What This Is, And Isn't.
§ 09 Significance — Why This Matters.
§ 10 Proposed Timeline — The Road Ahead.
§ 11 Selected References — The Shoulders.
Appendix A Complete Glossary — Ninety-Five Entries.

(Page numbers to be finalized on conversion; this preview demonstrates CGU formatting — 1" margins, 12pt Times, double spacing.)

§ 01 Problem Statement — Three Gaps, One System.

1.1 · The Clinical Gap

NSCLC is the leading cause of cancer mortality globally, accounting for approximately 2.1 million new diagnoses per year and 1.53 million deaths annually (GLOBOCAN 2022). Approximately 70% of patients are diagnosed at advanced stage (locally advanced or metastatic), when curative resection is no longer possible. Median overall survival at advanced stage ranges from 8 to 14 months, and five-year survival is approximately 35% for locally advanced disease and 7% for metastatic disease — a figure itself inflated by the subset of patients who carry actionable driver mutations (EGFR, ALK, KRAS G12C) and benefit from approved targeted therapies. For the majority who do not carry actionable mutations, median survival and five-year rates are substantially lower. Advances in targeted therapy have transformed outcomes for patients with those actionable drivers — but they represent a minority. The majority of NSCLC patients carry mutations that current drug design approaches cannot address. They have no approved targeted therapy.

Table 1 — Pipeline-Eligible Non-druggable NSCLC Mutations

Nine rows across the tumor-suppressor genes carry active-space electron/qubit estimates and are catalogued here as candidates satisfying the BQP-relevant classification criteria in principle: no approved targeted therapy, fermionic active space within quantum feasibility window, and no tractable classical solution at chemical accuracy. Candidacy is not the same as computability, and the two should not be conflated — a residue-level structural coverage audit following this table (see below) found that only three of these rows, plus TP53 Y220C, currently have a cited crystal structure that actually resolves the mutated residue, and only three of those four are usable once Y220C is excluded on druggability grounds. The full pipeline-eligible set spans eleven genes; the two additional genes — SETD2 and U2AF1 — appear only in the patient-burden analysis of Table 2 and have no active-space entry at all yet.

Gene	Mutation	Mutation Type	Protein Location	Active Electrons	Qubits (JW)	Primary Cancer Types	NSCLC Frequency	Key Data Sources
TP53	C275F	Structural (hydrophobic core)	DNA-binding domain, loop-sheet-helix	~44	~88	NSCLC, breast, colorectal †	<1% (anchor case)	COSMIC, IARC TP53 DB, PDB 2OCJ
TP53	R175H	Structural (global destabilization)	DNA-binding domain, β -sandwich core	~44	~88	NSCLC, ovarian, colorectal	~3–4%	Joerger & Fersht (2007), TCGA
TP53	G245S	Structural (L3 loop)	DNA-binding domain, L3 loop	~40	~80	NSCLC, head & neck	~1–2%	IARC TP53 DB, COSMIC
TP53	R249S	Structural (L3 loop)	DNA-binding domain, L3 loop	~40	~80	Hepatocellular, NSCLC	<1%	COSMIC, IARC TP53 DB
TP53	R282W	Structural (H2 helix)	DNA-binding domain, H2 helix	~42	~84	NSCLC, breast	~1%	Joerger & Fersht (2007), TCGA
KEAP1	G333C	LOF (Kelch domain)	Kelch β -propeller, NRF2 interface	~28	~56	NSCLC (adenocarcinoma)	~2–3%	TCGA, COSMIC, PDB 2FLU
KEAP1	R320Q	LOF (Kelch domain)	Kelch β -propeller, NRF2 interface	~28	~56	NSCLC (adenocarcinoma)	~1–2%	TCGA, COSMIC
CDKN2A	p16 LOF	LOF (deletion/inactivation)	CDK4/6 inhibitor interface	~24	~48	NSCLC, melanoma, pancreatic	~15–20%	TCGA, PDB 2A5E; [62] [63]
STK11	F354L	Kinase domain LOF	C-terminal non-catalytic region (unresolved)	~26	~52	NSCLC (adenocarcinoma)	~2–3%	TCGA, COSMIC — no PDB coverage; residue 354 falls outside 2WTK's resolved range (43-347), in LKB1's C-terminal region, which has no identifiable folded domain
STK11	D194N	Kinase domain LOF	Kinase catalytic domain	~26	~52	NSCLC (adenocarcinoma)	~1–2%	TCGA, COSMIC
NF1	LOF	LOF (truncating)	GTPase-activating domain	~30	~60	NSCLC, glioma	~5–7%	TCGA, COSMIC
RB1	LOF	LOF (deletion/truncating)	Pocket domain (E2F interface)	~32	~64	NSCLC (SCLC), retinoblastoma †	~3–5%	TCGA, COSMIC
SMARCA4	LOF	LOF (ATPase domain)	SWI/SNF ATPase domain	~28	~56	NSCLC, ovarian, medulloblastoma	~5–8%	TCGA, COSMIC
RBM10	LOF	LOF (splicing regulator)	RNA-binding domain	~22	~44	NSCLC (adenocarcinoma)	~5–7%	TCGA, COSMIC
ARID1A	LOF	LOF (chromatin remodeling)	ARID DNA-binding domain	~24	~48	NSCLC, endometrial, gastric †	~4–6%	TCGA, COSMIC

† Also observed in urological malignancies (urothelial carcinoma, renal cell carcinoma); SOLANGE pipeline directly applicable without architectural modification.

Structural coverage audit — a residue cited is not a residue resolved. Every row above names a mutation site and cites a PDB structure or a literature source; neither citation by itself establishes that the structure actually resolves the mutated residue, and this dissertation had not previously checked. Prompted by an external reviewer (M. Schmid) relaying a specialist's caution (A. Ray) that TP53-inactivating mutations are not all confined to well-resolved structural regions — some sit in intrinsically disordered regions (IDRs) invisible to X-ray or cryo-EM — every cited PDB entry in Table 1 was checked residue-by-residue against its actual resolved range. Two independent failures surfaced. KEAP1 R320Q : the Kelch domain's continuous backbone density begins at residue 322 in both cited structures (1U6D, 2FLU); R320 sits two residues before that boundary, inside a linker to the IVR domain independently documented in the literature as flexible and as the specific obstacle that has prevented a full-length Keap1 structure from being solved — not a marginal citation gap but a documented disordered region. STK11 F354L : the PDB code originally on record for this row, 2QK7, is not an LKB1 structure at all — it is a staphylococcal γ -hemolysin, an unrelated bacterial toxin. LKB1's actual solved structure (2WTK, Zeqiraj et al. 2009) resolves residues 43–347; F354 sits outside that range, in a C-terminal non-catalytic region the literature describes as having no

identifiable folded domain. KEAP1 LOF (the generic loss-of-function row, corresponding to the platform's demonstrated NGS finding of a frameshift at residue 196) fails on two independent grounds: residue 196 sits in the IVR, before the Kelch domain entirely, and a frameshift has no defined single mutant residue for an active space to be built around in the first place — the model compound used for this row (methanethiol, borrowed from the unrelated point mutation G333C) is a proxy standing in for a proxy. CDKN2A 's homozygous-deletion row is structurally inapplicable for the same reason as any deletion: there is no point residue to anchor an active space to. Three rows were independently re-verified as resolved and residue-accurate: TP53 C275F (2OCJ, residues 96–289; C275 sits centrally within that range; local active site 24e/48q, full site 44e/88q), KEAP1 G333C (1U6D/2FLU, residues 322–609; G333 sits well inside; local active site 104e/208q, full site 155e/310q), and STK11 D194N (2WTK, residues 43–347; D194 sits centrally within; local active site 76e/152q, full site 152e/304q). TP53 Y220C is separately excluded from the pipeline-eligible set on druggability grounds (§ below), leaving these three — one verified point mutation per core gene — as the only entries in this table with both a real structural anchor and a defensible non-druggable claim as of this writing.

Structural validity alone does not make a target reachable today. Of the three, only TP53 C275F 's footprint (44e/88q at the full site) sits within the range this dissertation's own hardware-feasibility argument supports for current-generation devices: IBM's demonstrated 32e/100q execution (arXiv:2603.08696) places 88 qubits within the same generation of machine, not a future one — the reasoning behind treating C275F as executable on IBM Heron r3 today (§ 05, Phase 3B). KEAP1 G333C exceeds that reach even at its reduced local site — 208 qubits is beyond Heron r3's own capacity (~156q), let alone the 310 qubits its full site requires. STK11 D194N 's local site (152q) sits almost exactly at that same ceiling; whether it is reachable is a hardware-capacity question this dissertation cannot resolve in the abstract, but its hardware_era designation in this pipeline is fault-tolerant (~2030+), not near-term, because qubit count alone does not establish that a circuit of that depth is executable at usable fidelity on a NISQ-era device. The honest statement, then, is not "three validated targets" without qualification: it is one target executable now (C275F), and two structurally valid targets whose execution is a statement about the 2030s, not about this dissertation's present hardware access.

This audit is reported here rather than silently corrected because the failure mode it caught — a citation that looks authoritative but does not withstand a residue-level check — is exactly the kind of error the evidence-discipline principles of § 06 exist to prevent, and finding one in the dissertation's own reference table is a more convincing demonstration of that discipline than never having looked. The same check was subsequently extended past Table 1's original four genes to all eleven genes underlying Table 2's patient-burden estimate, asking of each: is there a real, recurrent point-mutation hotspot (not a truncating event with no single residue to anchor), does a structure resolve it, and is it still undruggable. Table 1b reports the result.

Table 1b — SOLANGE-Eligible Point Mutations, Extended Audit

Result of extending the residue-level structural check across all eleven Table-2 genes, not only the original four. NF1, RB1, RBM10, and ARID1A were also checked and found to have no qualifying candidate: NF1 and RBM10 mutations are spread across the gene with no recurrent residue; RB1's one missense candidate (I680T) is a single case-report finding, not a validated hotspot; ARID1A's real missense hotspot cluster (residues ~1000–1119, structurally resolved by PDB 1RyU) is documented in hepatocellular, cutaneous, uterine, and gastric carcinoma, not NSCLC — included below anyway as ARID1A R1020S, explicitly flagged non-NSCLC, because it passes every gate except tissue relevance. U2AF1 S34F/S34Y is a real, NSCLC-recurrent hotspot (~81%/~19% split of U2AF1 mutations in lung adenocarcinoma) that fails structurally: no human structure resolves residue 34, which sits in an N-terminal zinc finger outside the one solved U2AF1 fragment (1JMT, residues 43–146) — it is not listed below. CDKN2A, KEAP1 LOF, KEAP1 R320Q, STK11 F354L, and ARID2 LOF remain excluded per the findings above this table.

#	Gene / Mutation	PDB	Active e ⁻ / Qubits (full site)	Hardware Era	NSCLC-relevant	Status
1	TP53 C275F	2OCJ	44e / 88q	current	yes — anchor	VERIFIED (measured, not size prior)
2	TP53 R175H	2OCJ (WT)	44e / 88q	current	yes	VERIFIED
3	TP53 G245S	2J1Y	40e / 80q	current	yes	VERIFIED
4	TP53 R249S	2BIO	40e / 80q	current	yes	VERIFIED
5	TP53 R282W	2J21	42e / 84q	current	yes	VERIFIED
6	KEAP1 G333C	1U6D	155e / 310q	fault-tolerant	yes	VERIFIED
7	STK11 D194N	2WTK	152e / 304q	fault-tolerant	yes	VERIFIED
8	SETD2 R1625C	5JJY	30e / 60q	current	yes	VERIFIED — most frequently mutated epigenetic modifier in LUAD (~9%)
9	ARID1A R1020S	1RYU	28e / 56q	current	no — hepatocellular / cutaneous / uterine / gastric	VERIFIED, structure-only inclusion
10	SMARCA4 R1192C	7VDT / 7VDV	28e / 56q	current	yes	PLAUSIBLE, NOT FULLY VERIFIED — RCSB direct confirmation pending

Rows 1–8 are structurally verified against the actual deposited PDB residue range, the same standard applied throughout this audit — not the size-only heuristic Table 1 opened with. Row 9 passes every gate except NSCLC relevance and is kept only because the structural question was asked independently of tissue; it is not part of this dissertation's patient-burden claim. Row 10's structure could not be independently confirmed in the working session that found it — both RCSB and EBI/PDBe were unreachable from that environment — and its status is provisional pending direct verification, distinguished here from the nine confirmed rows rather than folded into them.

Table 2 — Annual Patient Burden by Mutation: Non-druggable NSCLC in the Quantum Feasibility Window

Per-mutation patient counts derived from NSCLC frequency estimates applied to 2.12M new diagnoses/year (GLOBOCAN 2022). The naïve sum (2,877,000) is not a patient count: it double-counts patients, because a patient carrying several of these co-occurring mutations is tallied once in each mutation's row, so the sum exceeds the ~2.12M total annual NSCLC incidence. Removing that overlap — a ~55–62% reduction ($2,877,000 \times \sim 0.40 \approx 1.1\text{--}1.3\text{M}$) — yields the estimated number of unique patients carrying at least one non-druggable mutation. This ~1.1–1.3M is roughly 52–61% of all NSCLC patients, consistent with the ~60% non-druggable fraction stated in the Abstract. All mutations satisfy the three BQP-relevant classification criteria: no approved targeted therapy, fermionic active space within quantum feasibility window, and no tractable classical solution at chemical accuracy. This burden analysis spans eleven genes — the nine classified in Table 1 plus SETD2 and U2AF1, which contribute to the population estimate but await full active-space classification.

Gene / Mutation	NSCLC Frequency	Est. Patients/Year	Notes
TP53 C275F (anchor case)	~2%	~42,000	Structural mutant, loop-sheet-helix junction; no PDB entry; no clinical trial as of April 2026
TP53 R175H	~5%	~105,000	Global destabilization; most common TP53 hotspot; gain-of-function activity reported
TP53 G245S	~3%	~64,000	DNA-contact region; partial contact + structural destabilization
TP53 R249S	~2%	~42,000	L3 loop; associated with hepatocellular carcinoma; NSCLC co-driver
TP53 R282W	~2%	~42,000	Hydrophobic core; structural class; no targeted therapy
KEAP1 LOF (G333C, R320Q + others)	~17%	~360,000	Predominantly adenocarcinoma; no approved direct KEAP1 inhibitor
CDKN2A p16 LOF	~17%	~360,000	CDK4/6 inhibitors approved in breast cancer; not NSCLC-specific
STK11/LKB1 LOF	~20%	~424,000	Major driver of immunotherapy resistance; no direct STK11 therapy
NF1 LOF	~6%	~127,000	RAS-pathway; MEK inhibitors partially address; no direct NF1 therapy
RB1 LOF	~4%	~85,000	Predominantly SCLC component; no direct RB1 therapy
SMARCA4 LOF	~6%	~127,000	SWI/SNF complex; no approved targeted therapy
RBM10 LOF	~6%	~127,000	Splicing regulator; no approved targeted therapy
ARID1A LOF	~5%	~106,000	Chromatin remodeling; no approved targeted therapy
SETD2 LOF	~5%	~106,000	Histone methyltransferase; tumor suppressor; no approved targeted therapy
U2AF1 S34F/Y	~3%	~64,000	Splicing factor mutation; no approved targeted therapy
Naïve sum (pre-correction)	—	~2,877,000	Exceeds annual incidence due to co-occurring mutations in the same patient
Overlap-corrected total (~55–62% correction)	—	~1.1–1.3M unique patients/year	Unique patients carrying ≥ 1 non-druggable mutation within the quantum feasibility window

The magnitude of this unaddressed population reflects a structural limitation in the current drug-discovery pipeline: *de novo* resistance . Even among mutations that initially yield to targeted therapy — such as EGFR exon 19 deletions responding to osimertinib — resistance emerges within 12–18 months in the majority of patients. For the non-druggable class represented in Table 2, no first-line targeted therapy exists to trigger resistance in the first place. The pipeline has not simply failed these patients; it has never begun. A historical analogy clarifies the scale of the unsolved problem: HIV infection was once uniformly fatal within years of diagnosis, with no therapeutic options and no drug-discovery framework capable of attacking the virus's protease at the required precision. The development of combination antiretroviral therapy — grounded in computational structural biology and iterative molecular simulation — transformed HIV from a death sentence into a manageable chronic condition. The quantum simulation platform proposed in this dissertation is not a claim to replicate that transformation; it is a claim to establish the foundational information systems infrastructure — governed, reproducible, and Part-11-compliant — that makes iterative computational drug discovery possible for the non-druggable tumor-suppressor class, in the same way that reliable viral load measurement and provenance-tracked resistance profiling made antiretroviral therapy iteration possible for HIV.

Among the most challenging are TP53 structural DNA-interface mutants . Unlike cavity-creating mutants such as Y220C, which creates a druggable pocket, structural DNA-interface mutants such as p.Cys275Phe disrupt the DNA-binding interface and impair p53 transcriptional function. The Cys-to-Phe substitution introduces a bulky aromatic side chain near the loop-sheet-helix DNA-contact region of the p53 DNA-binding domain, likely perturbing local DNA-binding geometry rather than creating a therapeutically exploitable cavity. No approved drug specifically targets TP53 C275F, and no C275F-specific clinical trial appears to be registered on ClinicalTrials.gov as of April 2026.

The clinical urgency of this problem is not hypothetical. The patient NGS data anchoring this research (report MI25-0349 , Emek Medical Center, June 2025) identified TP53 p.Cys275Phe as a Tier I pathogenic variant at allele frequency 12.66% in a pulmonary adenocarcinoma case — alongside KEAP1 p.Cys196Leufs* and CDKN2A deletion, with no EGFR, ALK, KRAS, or RET detected. This patient had no approved targeted therapy available.

The structural inaccessibility of C275F compounds the clinical gap. As of April 2026, no crystallographic structure of the TP53 C275F mutant exists in the RCSB Protein Data Bank. For comparison, the Y220C mutant — which creates a druggable pocket — has been crystallised and deposited (PDB: 2VUK), enabling structure-based drug design. C275F has not: the Cys → Phe substitution destabilises the hydrophobic core without producing a tractable crystal form, and the resulting conformational heterogeneity prevents X-ray structure determination at useful resolution. The wild-type backbone shown in Figure 1.1 is therefore not a starting point for classical docking — it is the closest available approximation to an unknown structure. The absence of a C275F PDB entry is not a gap in the literature that will be filled by a better crystallography experiment; it is a signal that the mutation's structural consequences must be computed, not observed. The electronic structure of C275F at this active-space scale lies beyond the reach of exact classical methods (CCSD(T)/CASSCF, which become intractable past ~18 active electrons); quantum simulation is the targeted approach of this work for that regime — not a claim that it is the sole conceivable route, but the one this artifact is built to govern and orchestrate.

This dissertation is a consumer of PDB data, not a producer of it. The structural record above (PDB 2OCJ) is an input to the orchestration pipeline — parallel to the clinical NGS report MI25-0349 that enters the same pipeline at the patient data layer. The platform's first operation upon receiving a mutation target is to retrieve the relevant PDB entry, extract the coordinate data required to define the active-space boundary, and log the PDB identifier and retrieval timestamp in the provenance record. What the platform does with that input — active-space selection, Hamiltonian encoding, quantum simulation routing, error mitigation, and Part-11-compliant output generation — is the IT contribution. The structural biology is upstream of this architecture; the information systems work begins at the point of ingestion.

Why C275F and Not Another Mutation?

The selection of TP53 p.Cys275Phe (C275F) as the anchor case satisfies five independent criteria, each of which alone would justify its choice, and which jointly make it a uniquely suitable demonstration target for a quantum-simulation orchestration architecture.

Criterion 1 — Therapeutic orphan status. C275F has no approved targeted therapy and no registered clinical trial as of April 2026. The platform is designed precisely for the class of mutations where classical drug-discovery pipelines have failed. Mutations with existing therapies (e.g., TP53 Y220C with Rezatapopt in Phase II) are poor demonstration targets because the platform would be solving an already-solved problem. C275F is unambiguously in the unsolved class.

Criterion 2 — Absence of crystallographic data. No structure of the C275F mutant exists in the RCSB Protein Data Bank. The Cys → Phe substitution destabilizes the hydrophobic core of the DNA-binding domain, producing conformational heterogeneity that prevents X-ray crystallography at useful resolution. The electronic structure of C275F cannot be observed — it must be computed. This makes quantum simulation not one option among several but the primary path to chemical-accuracy characterization, which is precisely the condition under which the BQP-relevant classification framework applies.

Criterion 3 — Active-space size in the operationally BQP-relevant range. The full binding-site model requires approximately 44 active electrons, mapping to approximately 88 qubits under Jordan-Wigner encoding. This places the problem above the classical tractability threshold for coupled-cluster methods and within the demonstrated range of the SqDRIFT framework (IBM, March 2026: 32 electrons on 100 superconducting qubits). C275F sits precisely in the regime the platform is designed to govern — neither too small for classical methods nor too large for near-term quantum hardware.

Criterion 4 — Structural mutation class strengthens the quantum argument. As established in §1.2, C275F is a structural mutant: the perturbation is delocalized across the hydrophobic core of the DNA-binding domain rather than local to a single DNA-contacting residue. This delocalization produces correlated electron behavior across multiple residues simultaneously — the regime where classical DFT fails at chemical accuracy and where fermionic quantum simulation provides a principled advantage. A contact mutant such as R248W, by contrast,

involves a more localized perturbation that is in principle more tractable classically. C275F is therefore a harder problem in exactly the right way.

Criterion 5 — Clinical anchoring in real patient data. The mutation is sourced from a de-identified clinical NGS report (MI25-0349, Emek Medical Center, June 2025) identifying TP53 C275F as a Tier I pathogenic variant at 12.66% allele frequency in a pulmonary adenocarcinoma case with no actionable targeted driver detected. Using a real clinical case as the motivating instance grounds the artifact in an authentic translational problem, consistent with DSR's requirement that artifacts address real, significant organizational and technical problems rather than constructed ones.

The selection of a single mutation as the anchor case is a deliberate methodological choice, not a claim that C275F is uniquely important among all non-druggable tumor-suppressor mutations. The platform is hardware-agnostic and mutation-agnostic at the workflow layer: the same orchestration architecture, provenance schema, and compliance framework apply to any fermionic target satisfying the BQP-relevant classification criteria. C275F is the existence proof; analytical generalization to other non-druggable structural mutants follows from the architecture's design, not from additional case studies.

1.2 · The Computational Gap

Classical computational chemistry cannot reliably model the binding energetics of TP53 C275F at the required accuracy level. DFT (B3LYP) produces errors of $\pm 5\text{--}10$ kcal/mol against the ≤ 2 kcal/mol chemical accuracy threshold. CCSD(T) achieves gold-standard accuracy (± 1 kcal/mol) but scales as $O(N^7)$, becoming computationally infeasible at ~ 44 active electrons. Full Configuration Interaction (FCI) scales as $O(2^N)$ — reaching ~ 17 trillion operations at $N=44$, rendering it completely intractable.

The full binding pocket of TP53 C275F requires approximately 44 active electrons and 44 active orbitals. Jordan-Wigner encoding maps this to ~ 88 qubits — within the capacity of IBM Heron r3 (156 qubits). VQE with sqDRIFT Hamiltonian simulation provides polynomial circuit-depth scaling, placing this problem within BQP under the standard quantum-chemistry framing (Abrams & Lloyd, 1999; Aspuru-Guzik et al., 2005) and making the active space tractable on NISQ-era hardware.

The structural — rather than contact — nature of C275F further strengthens the quantum simulation argument on three independent grounds. First, structural mutations produce delocalized electronic perturbations across the entire active site: the Cys $\rightarrow$ Phe substitution reorganizes the hydrophobic core of the DBD, creating correlated electron behavior across multiple residues simultaneously. Classical DFT methods, which treat exchange-correlation locally, cannot capture this non-local electronic reorganization at chemical accuracy — precisely the regime where SqDRIFT's polynomial-time fermionic simulation provides advantage. Second, the absence of a C275F crystal structure in the RCSB PDB is itself a direct consequence of structural instability: conformational heterogeneity induced by the mutation prevents X-ray structure determination at useful resolution, confirming that the electronic consequences of C275F must be computed rather than observed. This makes quantum simulation not merely advantageous but the only viable path to chemical-accuracy characterization. Third, the 44-electron active space required for Phase 3B is scientifically justified by the structural classification: unlike contact mutations — where the perturbation is local to a single DNA-contacting residue — structural mutations require inclusion of the broader electronic environment encompassing the Zn^{2+} coordination shell and adjacent hydrophobic core residues. The two-level active space framework (24 electrons local, 44 electrons full) is therefore not an arbitrary design choice but a direct consequence of C275F's structural mutation mechanism.

This dissertation does not claim a formal complexity-theoretic separation between classical and quantum solvability for this problem; such separations remain open in general. The claim is the weaker but defensible one — that C275F sits in the empirically quantum-advantaged regime as defined by the current state of both classical and quantum methods circa 2026. Advanced classical approximations are explicitly acknowledged: tensor-network methods (DMRG) are highly effective for one-dimensional and quasi-1D systems but degrade on three-dimensional, pocket-like active spaces with strong correlation across non-local orbitals; neural quantum states (Carleo & Troyer,

2017) have shown promising results on lattice models but have not yet demonstrated chemical accuracy on fermionic systems at the ~44-electron metalloprotein scale required here. The platform's BQP-relevant framework remains valid until — and only until — a classical method demonstrates chemical accuracy at this active-space boundary within tractable resources, at which point the architecture's hardware-agnostic abstraction would simply route the workload to that classical backend.

Two complexity-theoretic subtleties are stated directly, since they are easy to overstate. First, finding the ground state of a general local Hamiltonian is QMA-hard — the quantum analogue of NP-hardness — so no polynomial-time quantum algorithm is guaranteed in the worst case. The containment of a specific instance in BQP is conditional on preparing an initial state with sufficient overlap on the true ground state (the overlap problem), and constructing such a state is itself hard in general. This dissertation therefore makes no claim of guaranteed exponential quantum advantage. Second, and directly on point, Lee et al. (2023) — Evaluating the evidence for exponential quantum advantage in ground-state quantum chemistry — argue that for many ground-state chemistry problems the initial-overlap requirement and the strength of classical heuristics together erode the case for exponential speedup, concluding that evidence for exponential quantum advantage across chemical space has yet to be found. That critique is accepted here as correct in the general case; rather than undermining this work it delimits it. The claim advanced is not exponential advantage in principle but a practical quantum algorithm in the specific operational regime of the active spaces selected — where Hartree–Fock/CASSCF reference states supply adequate initial overlap and the best classical methods do not reach chemical accuracy at tractable cost — with the dissertation's contribution being the governance and orchestration architecture around that algorithm, not a proof of quantum supremacy for chemistry.

1.2.1 · Active Space Definition — Two-Level Framework

This subsection defines a system requirement for the orchestration architecture, not a chemistry claim. For the workflow to handle C275F end-to-end, it must commit to a precise active-space boundary; that boundary determines qubit count, telemetry volume, and audit-trail granularity downstream. The chemistry is anchoring detail; the IT contribution is the workflow's ability to execute reproducibly at either boundary. Two levels are defined, each corresponding to a distinct architectural validation phase.

Local Binding Site · ~24 Active Electrons

Encompasses the immediate chemical environment of the mutation: the F275 aromatic π -system (6 electrons), adjacent residues C277 and A276 (~10 electrons), and relevant bonding orbitals between them (~8 electrons). This level is sufficient for quantum-accurate drug candidate identification — characterizing the altered binding pocket geometry, evaluating candidate molecule interactions, and comparing wild-type versus mutant electronic structure. Requiring approximately 48 qubits under Jordan-Wigner encoding, this active space is directly feasible on IBM Heron r3 (156 qubits) and is enabled by the SqDRIFT benchmark (IBM Research, 2026; arXiv:2603.08696: 32 electrons demonstrated on 100 qubits).

Full Binding Site · ~44 Active Electrons

Extends the local site to incorporate the complete biologically active environment: the Zn^{2+} coordination shell and its four ligands (C176, H179, C238, C242), plus the DNA-guanine interface and phosphate backbone (~12–20 additional electrons). This level is required for full functional validation — confirming whether a drug candidate restores p53 DNA-binding activity. Jordan-Wigner encoding maps ~44 electrons to ~88 qubits, placing this within IBM Heron r3 capacity but at the frontier of demonstrated capability.

System	Active Electrons	Qubits Required	Feasibility
C_{13}Cl_2 (IBM Research 2026 benchmark)	32	~100	✓ Demonstrated
C275F — Local binding site	~24	~48	✓ Feasible today
C275F — Full binding site	~44	~88	► Near-term frontier

1.3 · The Systems Gap

Even if quantum hardware can simulate TP53 C275F, a deeper IT problem exists: no integrated system connects the clinical problem to a quantum simulation result in a governed, reproducible, stakeholder-accessible way. Quantum platforms are vendor-specific. HPC tools run in isolated environments. AI rescoring models require separate pipelines. No orchestration layer connects them. No compliance framework governs the output. No visualization layer communicates results to the oncologists, pharma researchers, and hospital IT teams who need to act on them.

1.4 · Hardware Modality Anchor

The modality choice is an IT decision, not a physics ranking. The orchestration architecture must commit to a single hardware anchor whose telemetry is well-specified, whose calibration epochs are exposed, and whose compiler stack supports the audit-trail framework of §06.iii — otherwise the platform cannot satisfy 21 CFR Part 11 reproducibility requirements at proof-of-concept. The four-modality survey below is the engineering rationale for that commitment, not a scientific judgement about which physical substrate is "best."

The orchestration architecture is hardware-agnostic at the workflow level; vendor-specific at the compilation and error-mitigation layer. IBM Heron, IonQ, neutral-atom systems, and HPC backends are all addressable through the same workflow abstraction, but compilation strategy (SWAP overhead on superconducting; all-to-all on trapped-ion; reconfigurable geometry on neutral atoms) and error-mitigation primitives (ZNE tuned for superconducting noise; different dominant error channels elsewhere) cannot be made fully portable. The contribution is to minimize this leak through a documented translation layer, not to eliminate it. The proof-of-concept execution, however, must anchor on a single quantum modality. The four commercially accessible paradigms in 2026 — superconducting circuits, trapped ions, neutral atoms, and photonic systems — differ materially in fidelity, connectivity, gate speed, and software maturity. The choice of anchor must be derived from the specific computational demands of this dissertation's target problem (gate-based ground-state estimation for fermionic Hamiltonians at 20–32 electron active spaces, executed as sample-based quantum diagonalization — quantum sampling of the dominant electronic configurations followed by classical diagonalization over the sampled subspace), not from generic benchmarks.

Eight criteria are dispositive for the present problem class: C1 active-space tractability, C2 circuit-depth fidelity, C3 wall-clock execution speed, C4 SqDRIFT compiler support, C5 HPC integration, C6 audit-trail telemetry, C7 empirical precedent at the target scale, and C8 timeline stability through the dissertation period.

Criterion	Superconducting	Trapped Ion	Neutral Atom	Photonic
C1 · Active-space scale	Strong	Adequate	Strong	Limited
C2 · Circuit-depth fidelity	Adequate	Strong	Adequate	Limited
C3 · Gate speed / wall-clock	Strong (ns)	Limited (ms)	Limited (μ s–ms)	Strong
C4 · SqDRIFT compiler	Strong (Qiskit)	Adequate	Emerging	Limited
C5 · HPC integration	Strong	Adequate	Emerging	Limited
C6 · Audit-trail telemetry	Strong	Adequate	Emerging	Limited
C7 · Precedent at target scale	Strong (SqDRIFT 2026)	Adequate	Adequate (different alg.)	Limited
C8 · Timeline stability	Strong	Strong	Adequate	Adequate

Decision: superconducting qubits are selected as the anchoring modality. The choice is the best-supported at submission for three converging reasons. First, the SqDRIFT framework — on which this dissertation's time-compression argument rests — was demonstrated in 2026 on 100 superconducting qubits recovering the ground-state energy of a 32-electron active space of $C_{13}Cl_2$. This is the only published benchmark in the world that matches the active-space size targeted here. Second, Qiskit Runtime provides production-grade primitives for fermionic Hamiltonian construction, sample-state preparation, classical diagonalization, and per-circuit telemetry sufficient to populate the audit-trail framework specified in §06. Third, SqDRIFT is fundamentally a sampling procedure — the quantum processor draws many

independent samples of the dominant electronic configurations, which a classical solver then diagonalizes — so its time-to-solution is dominated by the product of per-circuit execution time and the number of sampling circuits and shots. Nanosecond gates make this product tractable; millisecond gates would inflate it by three orders of magnitude.

The choice is anchored, not absolute. A dual-modality framing (superconducting + neutral atoms) would dilute the central thesis without empirical symmetry: no SqDRIFT-class workload has been demonstrated on neutral atoms at the relevant scale. The pipeline architecture is described abstractly enough that re-anchoring on a different modality, should the empirical landscape shift before defense, is a defined exercise in future work — not a re-derivation of the core results. Two modality-specific properties propagate forward: connectivity-induced SWAP overhead (treated as a sensitivity parameter in circuit-depth estimates) and calibration drift (treated as a mandatory metadata field in the simulation provenance record — see §06.iii).

1.4.1 · Post-April 2026 Hybrid Landscape

The modality argument above was framed against the state of the field at submission. On April 14, 2026, Nvidia released Ising — a 35-billion-parameter vision-language model for agentic qubit calibration paired with a 3D convolutional neural network that acts as a fast pre-decoder for real-time quantum error correction, resolving the common syndrome patterns and escalating the remainder to a conventional global decoder (e.g., pyMatching), reportedly 2.5× faster and 3× more accurate than the pyMatching-only baseline — alongside the NVQLink GPU ↔ QPU interconnect and continued investment in the CUDA-Q programming model. Nvidia's stated position, confirmed by CEO Jensen Huang, is that the company will not build a quantum processing unit; it intends to occupy the classical, AI, and HPC substrate sitting beneath any vendor's QPU. This release does not weaken the modality argument made above; it confirms its structural premise. The QC·AI·HPC stack is bifurcating into a quantum-kernel layer (QPU + algorithm) and a hybrid-substrate layer (classical pre/post-processing, AI calibration, AI-decoded QEC). IBM Heron r3 + SqDRIFT remains this dissertation's quantum-kernel anchor — and IBM is notably absent from Nvidia's announced NVQLink partner roster (Atom Computing, IonQ, IQM, Q-CTRL, Academia Sinica, EeroQ), continuing its own Quantum-Centric Supercomputing path on Qiskit — independently validated at protein scale on March 26, 2026, when Cleveland Clinic and IBM published the first electronic-structure simulation of the 303-atom Trp-cage miniprotein via Embedded Wave Function with Sample-based Quantum Diagonalization (SQD) on Heron r2. The orchestration architecture defended in §06.ii is therefore positioned to sit across this bifurcation: treating CUDA-Q, NVQLink, and Ising as commodity substrate where they apply, and Qiskit Runtime where IBM's path is preferred — rather than committing the dissertation to either side of the emerging vendor split.



[Figure — rendered from source]

1.4.2 · Trapped-Ion Logical-Qubit Landscape

On March 26, 2026, Quantinuum reported a demonstration on its Helios trapped-ion processor of quantum computations with up to 94 logical qubits using error-detection codes (and up to 48 logical qubits using concatenated error-correction codes), achieving "beyond break-even" performance — encoded qubits producing lower error rates than unencoded physical hardware operations. This is a notable milestone for the trapped-ion modality and sharpens — but does not invalidate — the superconducting modality argument made in §1.4.

A common misconception in evaluating quantum hardware for computational chemistry is to treat coherence time as the primary figure of merit. Trapped-ion systems exhibit coherence times measured in seconds — occasionally minutes in specialized configurations — while superconducting transmon systems such as IBM's Heron operate with coherence windows of approximately 100–300 microseconds. Taken in isolation, this represents a four-orders-of-magnitude advantage for trapped ions.

The operative metric, however, is not absolute coherence but the ratio between coherence time and gate time, which determines the effective circuit depth available before decoherence dominates. Trapped-ion systems combine multi-second coherence with microsecond-scale gate operations, yielding on the order of 10^6 operations per coherence window. Superconducting systems combine ~200-microsecond coherence with nanosecond-scale gates, yielding on the order of 10^4 operations per window. By this measure, trapped ions retain a meaningful advantage for deep circuits.

This advantage does not translate uniformly across application domains. Drug-discovery workloads — particularly the simulation of mutation-specific electronic structure across libraries of candidate ligands and conformational variants — are not characterized by a single deep circuit, but by a high volume of moderate-depth circuits executed across a broad parameter space. SqDRIFT — the algorithm anchoring this dissertation's empirical claim — is a sample-based quantum diagonalization method: the quantum processor samples the high-weight electronic configurations that span the relevant subspace, and a classical processor diagonalizes the Hamiltonian projected onto those samples. Reaching a converged energy requires a large number of independent sampling circuits and shots, so its time-to-solution is determined by the product of per-circuit execution time $\times$ sample count. This is the canonical throughput-bound regime: aggregate throughput, not per-circuit depth, is the binding metric. Under this regime, the nanosecond gate times of superconducting hardware confer a decisive practical advantage despite the shorter coherence window per circuit.

A March 2026 Nature report from Quantinuum on its Helios processor provides a concrete, citable data point for this throughput argument [55]. Helios is a 98-qubit trapped-ion QCCD device using ^{137}Ba + hyperfine qubits with all-to-all connectivity, achieving single-qubit infidelity of 2.5×10^{-5} , two-qubit gate infidelity of 7.9×10^{-4} (99.92% fidelity), and SPAM infidelity of 3.3×10^{-4} — fidelities that substantially exceed those reported for IBM Heron r3, and sufficient to demonstrate random circuit sampling beyond classical simulability. These figures confirm, rather than undermine, the coherence-versus-throughput distinction drawn above: the same report measures approximately 55 ms of wall-clock time per depth-1 entangling layer across the full 98-qubit register. A SqDRIFT sampling run, whose cost scales with the product of per-circuit execution time and sample count rather than with single-shot fidelity, would accumulate this per-layer latency across the thousands of sampling circuit instances the method requires. At Heron r3's nanosecond-scale gate times, the equivalent sampling budget completes in a small fraction of the time implied by Helios's millisecond-scale layer latency, even though Helios's per-gate error rates are lower. The Quantinuum result therefore strengthens rather than threatens the case made here: it is direct experimental confirmation that high coherence and high throughput are independent axes of merit, and that workload class — not platform reputation — determines which axis dominates.

The architectural choice of IBM Heron r3 as the hardware substrate for this dissertation's QC·AI·HPC pipeline therefore reflects a workload-aligned optimization, not a claim of categorical superiority. Quantinuum's logical-qubit demonstration does not weaken the modality argument made in §1.4; it sharpens its scope. Trapped-ion platforms remain well-suited to applications demanding maximal circuit depth per execution — fault-tolerance research, logical-qubit benchmarking, and certain classes of long-time-evolution simulation — domains this dissertation explicitly does not target. Drug discovery is throughput-dominated, and superconducting hardware is presently the modality whose operational profile aligns with that demand. The choice of the r3 revision specifically is a fidelity choice rather than a capacity one: r3 carries the same 156-qubit heavy-hex layout as r2, but delivers IBM's best Heron coherence to date ($\sim 350 \mu\text{s}$ T_2) with improved two-qubit gate and readout fidelity (IBM Quantum, 2025–2026 [58, 59]). Because C275F's ~ 88 -qubit active site already fits within r2's qubit count, the binding-relevant constraint is not more qubits but lower error — the reduced gate/readout error of r3 is what narrows the error-mitigation burden required to approach chemical accuracy on an 88-qubit circuit.

1.4.3 · Microsoft–Quantinuum Fault-Tolerant Milestone

A preceding and complementary milestone to the Helios results discussed above was published by Microsoft and Quantinuum in Nature under the title "Improved quantum processor logical error rates via correction and detection" [54]. Running on Quantinuum's H2 trapped-ion processor with Microsoft's qubit-virtualization software stack, the joint team demonstrated an 800-fold reduction in logical error rate relative to the physical qubit baseline — the largest recorded gap between physical and logical error rates reported in the peer-reviewed literature at the time. The system executed 14,000 individual circuit instances without a single observed error and scaled from 4 to 12 highly

reliable logical qubits using active syndrome extraction with real-time classical co-processor correction, without disturbing the quantum data during error diagnosis.

The significance for this dissertation is threefold. First, it constitutes independent peer-reviewed validation that the field has moved past the NISQ era into what Microsoft terms "resilient quantum computing" — the regime in which logical qubits outperform physical qubits, the foundational requirement for fault-tolerant quantum chemistry. Second, it demonstrates that the QEC machinery required for Phase 3B of the SOLANGE pipeline (error mitigation on IBM Heron r3, escalating to full fault-tolerance as hardware matures) is not a distant theoretical aspiration but an experimentally demonstrated trajectory. Third, the end-to-end chemistry simulation paired with AI and HPC supercomputers in the same demonstration is structurally identical to the QC·AI·HPC orchestration architecture defended in §06 — confirming the architectural pattern independently. This milestone does not change the hardware choice: Microsoft's qubit-virtualization stack runs on Quantinuum's trapped-ion hardware, not on IBM's superconducting backend, and the throughput argument of §1.4.2 continues to apply. It does, however, materially strengthen the scientific timeline argument for Phase 3B feasibility.

1.4.4 · Neutral-Atom Landscape (Infleqtion)

The neutral-atom modality reached a commercial-maturity milestone on February 17, 2026, when Infleqtion (formerly ColdQuanta) became the first neutral-atom quantum company to list publicly, trading on the NYSE under the ticker INFQ following a business combination with Churchill Capital Corp X. The listing raised over \$550 million in gross proceeds, and Infleqtion's systems are already deployed with the U.S. Department of War, NASA, and the U.K. government, alongside a joint demonstration with Nvidia of a materials-science application using logical qubits [56]. This commercial trajectory does not, on its own, alter the modality argument made in §1.4; it raises the question of whether the "Emerging" ratings assigned to neutral atoms in the C4-C6 rows of the comparison table still hold, and the answer requires separating two distinct claims that are easy to conflate.

The first claim is scale. Neutral-atom arrays, built from optically trapped alkali atoms with no fabrication-induced variation between qubits, scale favorably in qubit count: arrays exceeding several thousand trapped atoms have been demonstrated, and Infleqtion's own roadmap targets 100 logical qubits by 2028 on its Sqale platform. Google's public framing of its own neutral-atom research program captures this precisely: superconducting processors are easier to scale in the time dimension (circuit depth), while neutral atoms are easier to scale in the space dimension (qubit count). This is the basis for the "Strong" C1 rating neutral atoms already carry in the §1.4 table, and Infleqtion's commercial momentum is consistent with it.

The second claim is throughput, and this is where the modality does not transfer to the present problem class. Neutral-atom gate operations require physical qubit shuttling - optical tweezers must move atoms into proximity before an entangling gate can execute - and published two-qubit gate times sit near 400 nanoseconds, roughly an order of magnitude slower than superconducting gate times, with end-to-end computation reported at one-hundredth to one-thousandth the speed of superconducting systems. SqDRIFT, as established in §1.4.2, is a throughput-bound algorithm: its time-to-solution scales with the product of per-circuit execution time and sample count, not with active-space size in isolation. A modality that scales favorably in qubit count but unfavorably in per-circuit gate speed is solving the wrong half of the SqDRIFT cost function for this dissertation's purposes.

This sharpens, rather than duplicates, the trapped-ion analysis of §1.4.2. Trapped-ion systems trade slow gates for exceptionally long coherence times, preserving a real advantage for deep-circuit, fault-tolerance-oriented work. Neutral atoms do not carry an equivalent compensating advantage for the throughput-bound regime: coherence times are good but not in the trapped-ion class, and no SqDRIFT-class benchmark - or comparable sample-based quantum diagonalization result - has been published on neutral-atom hardware at the active-space scale this dissertation targets, which is why C7 remains rated "Adequate (different alg.)" rather than "Strong." Infleqtion's software stack, Superstaq, is a cross-hardware compilation and optimization layer; nothing in its public description indicates a Part 11-aligned provenance

or audit-trail framework comparable to the P1-P9 schema specified in §06.iii, leaving C5 and C6 correctly rated "Emerging" rather than revised upward.

Infleqion's public listing is genuine evidence that the neutral-atom modality is commercially de-risking - it is not evidence that the modality is better matched to this dissertation's throughput-bound workload. The decision recorded in §1.4 stands: superconducting qubits remain the best-supported anchor for a SqDRIFT-class pipeline at submission, and neutral atoms are correctly positioned in §1.4's table as a modality whose comparative advantage lies in a different axis of scaling than the one this dissertation's algorithm is bound by.

§ 02 Research Questions — Four Questions.

Can an orchestration layer that is bifurcation-aware (workflow-portable across IBM-native and Nvidia-substrate quantum stacks) and Part-11-governed (emitting nine-element provenance records, P1–P9, as first-class artifacts) unify quantum computing, HPC, and AI backends into a single reproducible workflow for non-druggable NSCLC tumor-suppressor mutation simulation?

Does the orchestration architecture correctly execute a VQE workflow end-to-end — from NGS input through active space selection, qubit encoding, Hamiltonian simulation, error mitigation, and AI rescoring — producing outputs consistent with quantum chemistry benchmarks on a simulator backend?

What design patterns and controls are necessary and sufficient to bring quantum-computed molecular binding data into compliance with FDA 21 CFR Part 11 electronic records requirements?

How should complex molecular structural information and quantum simulation results be presented to cross-disciplinary stakeholders — oncologists, pharma R&D, hospital IT — who lack quantum computing expertise?

§ 03 Literature Review & Research Gaps — What Exists, What's Missing.

3.1 · The Classical → Quantum Lineage and the Tractability Ceiling

The theoretical case for simulating chemistry on quantum hardware forms an unbroken lineage. Feynman (1982) first argued that quantum systems are intractable to classical simulation and that a computer built of quantum elements could simulate them efficiently; Lloyd (1996) made this constructive, proving that a universal quantum computer can efficiently simulate any local quantum system through Trotterized time evolution. Abrams and Lloyd (1999) then supplied the phase-estimation algorithm that converts simulated dynamics into ground-state eigenvalues — the formal bridge from "simulate the system" to "extract its energy" — and Aspuru-Guzik et al. (2005) mapped that machinery onto real molecules. McArdle et al. (2020) review the modern near-term descendants of this lineage. This continuum, from physical principle to executable molecular algorithm, is the theoretical spine on which the present architecture rests.

The lineage matters because the classical alternative has a definite ceiling. Exact correlation methods — full configuration interaction and the CASSCF/CASCI family — scale combinatorially with active-space size (Helgaker, Jørgensen, & Olsen, 2000; Levine et al., 2021), placing a practical wall near roughly eighteen active electrons. Reiher et al. (2017) quantify the crossover concretely, showing that classically intractable electronic-structure problems become reachable on fault-tolerant quantum hardware. The C275F target of this dissertation (~88 qubits under Jordan–Wigner encoding) sits decisively beyond that classical wall — which is precisely why a quantum route is the appropriate one.

On the near-term algorithmic side, Peruzzo et al. (2014) introduced the variational quantum eigensolver (VQE) as the hybrid loop pairing a quantum expectation-value estimate with a classical optimizer; McClean et al. (2016) established its theoretical framework; and Preskill (2018) named the NISQ regime that makes such hybrid methods the leading near-term strategy. Kandala et al. (2017) demonstrated hardware-efficient VQE on a superconducting processor, Bauer et al. (2020) surveyed the broader resource landscape, and Tilly et al. (2022) consolidated VQE methods and best practices. The randomized-compilation branch — Campbell's (2019) qDRIFT and its stochastic-sampling refinements (Kiss et al., 2023) — supplies the stochastic state-sampling basis for the sample-based quantum diagonalization (SqDRIFT) step used in Phase 3B. None of this work addresses specific oncology mutation targets, clinical translation, or the systems infrastructure required to operationalize quantum chemistry results.

3.2 · TP53 Mutation Landscape

Joerger and Fersht (2016) provide the definitive structural analysis of TP53 mutant classes, distinguishing cavity-creating mutants (Y220C) from DNA-interface structural mutants (C275F) and direct contact mutants (R248W). No published work applies quantum simulation to TP53 DNA-interface structural mutants such as C275F. No computational pathway exists in the literature for C275F drug discovery.

Figure 3.1 documents the platform's input diversity, not its scientific scope. Each of the five PDB entries is a data source the orchestration architecture must consume: a machine-readable record whose coordinate geometry determines active-space size, qubit demand, and backend routing for that specific job. The five targets differ significantly across all three dimensions — from the small ankyrin-repeat scaffold of CDKN2A p16 (156 residues, PDB 2A5E) to the Kelch β -propeller of KEAP1 (302 residues, PDB 2FLU) to the zinc-coordinated DNA-binding domain of TP53 C275F (~88 qubits under Jordan–Wigner encoding). Managing that heterogeneity through a single governed pipeline — same provenance schema, same compliance output, same audit trail — without manual re-derivation per target is the information systems design problem this dissertation solves. The platform consumes these five records; it does not produce them.

3.3 · Quantum Computing in Drug Discovery

Cao et al. (2019) reviewed quantum computing applications in quantum chemistry. Santagati et al. (2024) demonstrated quantum simulation of small drug-like molecules on real hardware. No published work applies these methods to non-druggable NSCLC tumor-suppressor structural mutants or integrates them into a governed clinical workflow.

3.4 · Hybrid Orchestration Platforms

Qiskit Runtime (IBM, 2021) and Orquestra (Zapata Computing, 2020) represented the first wave of single-vendor quantum orchestration. The post-March-2026 landscape is materially different: Nvidia's CUDA-Q programming model and the NVQLink GPU ↔ QPU interconnect (announced October 2025, expanded April 14, 2026 with the Ising AI models) constitute a cross-vendor, cross-paradigm orchestration substrate, and hybrid drug-discovery pipelines on it have been demonstrated — IonQ with AstraZeneca, AWS, and Nvidia on Suzuki–Miyaura coupling (June 2025) and Quantinuum with Nvidia on ADAPT-GQE applied to imipramine. Both precedents target generic small-molecule chemistry, not oncology mutation simulation, and neither is designed against a regulatory-records framework. What remains absent is therefore not a generic cross-vendor orchestration platform — that gap has begun to close — but a cross-vendor orchestration architecture specifically targeting oncology tumor-suppressor mutation simulation, with Part-11-grade provenance machinery sufficient to bring quantum-derived binding data into a regulated drug-development workflow. This is the narrower, defensible gap that §06.ii and §06.iii jointly close.

The post-April-2026 landscape was further shaped on March 26, 2026 by Cleveland Clinic and IBM's published quantum-centric supercomputing (QCSC) workflow — the first protein-scale electronic-structure simulation under that paradigm, whose architectural vision of tightly coupled HPC and quantum resources was set out by Alexeev et al. (2024) — which decomposed the 303-atom Trp-cage miniprotein into clusters and routed the most entangled (6–33 MOs each) to IBM Heron r2 via Embedded Wave Function with Sample-based Quantum Diagonalization (SQD). The Cleveland Clinic precedent is a methodological neighbor to this dissertation's sqDRIFT anchor (§1.4.1) and a structural neighbor to the bifurcation-aware orchestration claim (§06.ii), but it targets generic biomolecular discovery, not oncology-specific non-druggable mutations, and operates without a clinical-regulatory governance layer (no Part 11, no P1–P9 provenance record, no ML-decoder lineage capture). The narrower gap closed by this dissertation — oncology-mutation-specific orchestration with workflow-embedded compliance machinery — is therefore unaffected by the Cleveland Clinic precedent; if anything, it is sharpened.

3.5 · Regulatory Frameworks for Quantum Data

FDA 21 CFR Part 11 governs electronic records in regulated industries. Existing guidance covers classical computational models but does not address quantum-computed outputs. The April 14, 2026 release of Nvidia's Ising — a 3D convolutional neural network operating as a learned pre-decoder that resolves common syndromes at runtime and escalates the remainder to a conventional global decoder — sharpens the regulatory gap rather than closing it: a regulated drug-discovery pipeline must now produce Part-11-grade records not only of probabilistic quantum sampling (statistical estimators over noisy hardware) but also of a learned probabilistic pre-decoder in the QEC loop (and of the global decoder it defers to), whose model identifier, weight checksum, training corpus, and inference-time configuration become first-class compliance artifacts. The pattern is not unprecedented — FDA's evolving Good Machine Learning Practice (GMLP) guidance already requires model provenance for Software-as-a-Medical-Device — but it has not been extended to the QEC layer of a quantum simulation pipeline. No published framework applies Part 11 controls to quantum-derived binding data, and none addresses ML-decoded quantum error correction in a regulated workflow. This is the regulatory gap that §06.iii closes; the eight-element provenance record (P1–P8) is extended with a ninth element (P9) to capture ML-decoder lineage explicitly.

3.6 · Enterprise Architecture and Scientific-Data Provenance

The discipline that supplies a vocabulary for designing and governing complex systems-of-systems is enterprise architecture, established by Zachman (1987) and operationalized today through the TOGAF Standard and its Architecture Development Method (The Open Group, 2022), with ArchiMate (The Open Group, 2022) as the companion modeling language and Lankhorst (2017) as the standard scholarly treatment. These frameworks have been applied to scientific computing — Pierantoni et al. (2023) derive a reference architecture for science gateways spanning clouds, grids, and HPC — but the literature stops at the classical substrate: no published work frames a hybrid classical → quantum scientific pipeline as a governed enterprise architecture (as of June 2026; see the novelty-verification note in §06). Conversely, the quantum-chemistry literature reviewed in §3.1 is essentially silent on architecture governance.

The governance layer that makes such an architecture auditable rests on mature provenance and reproducibility standards. The FAIR principles (Wilkinson et al., 2016) define the findability, accessibility, interoperability, and reusability criteria a records layer must satisfy; the W3C PROV ontology (W3C, 2013) supplies an interoperable data model for serializing a provenance chain; and the reproducibility programs of Sandve et al. (2013) and Stodden et al. (2016) specify the operational and policy disciplines for trustworthy computational results. These standards, however, were written for conventional computation. None has been extended to a probabilistic, sampling-based quantum pipeline operating under regulatory-grade controls (21 CFR Part 11). The P1–P9 provenance schema of this dissertation is the instrument that conjugates these two otherwise-disconnected literatures — enterprise-architecture governance and scientific-data provenance — onto the classical → quantum continuum.

§ 04 Theoretical Framework — The Intellectual Scaffolding.

Design Science Research (DSR)

Hevner et al. (2004). The primary research method: building an artifact and evaluating it against defined criteria. The artifact is the contribution — its completeness, coherence, novelty, and utility are the evaluation criteria, not the scale of its deployment.

BQP-Relevant Quantum Chemistry

Bernstein & Vazirani (1997); Abrams & Lloyd (1999); Aspuru-Guzik et al. (2005). The theoretical grounding for placing the fermionic Hamiltonian eigenvalue problem at C275F's active-space boundary within BQP under standard quantum-chemistry framing. Establishes why quantum hardware is the appropriate substrate — under current classical state-of-the-art — without claiming a formal classical-quantum separation that the complexity literature has not proven for this problem class.

Translational Informatics

Sarkar (2010). The discipline of designing information systems that bridge clinical data, scientific computation, and actionable medical knowledge. Frames the platform as an IT artifact serving cross-disciplinary stakeholders rather than a physics instrument serving specialists.


Enterprise Architecture & the Classical–Quantum Continuum

The Open Group (2022); Lankhorst (2017). SOLANGE is not merely a quantum-chemistry pipeline but an enterprise architecture in the formal sense, and it is described here through TOGAF — the most widely adopted architecture framework. TOGAF contributes two instruments. First, a disciplined vocabulary that decomposes a complex socio-technical system into four architecture domains. Second, the Architecture Development Method (ADM), which models architecture not as a static blueprint but as a governed, iterative migration — the precise conceptual frame for SOLANGE's evolution from a classical proxy to quantum hardware.

TOGAF Domain	Framework Definition	SOLANGE Instantiation	Provenance
Business	Strategy, drivers, stakeholders, and the problem the system exists to solve.	Non-druggable NSCLC tumor-suppressor targets (TP53, KEAP1, CDKN2A, STK11) — mutations with no small-molecule binding pocket.	—
Data	Logical and physical data assets and their governance.	Jordan–Wigner-encoded molecular Hamiltonians, CASSCF reference datasets, and the immutable P1–P9 provenance records (21 CFR §11.10(e)).	P1 · P2 · P5
Application	Applications, their behavior, and their interactions.	CASSCF/PySCF (classical reference), VQE on PennyLane (live Phase 3A), and SqDRIFT sample-based quantum diagonalization (Phase 3B).	P1 · P6
Technology	Hardware, networks, and the platform infrastructure beneath everything above.	The classical → quantum bridge itself: academic HPC (NSF ACCESS-class allocations) on one side, IBM Heron r3 on the other.	P3 · P4

The ADM frames Phase 3A → 3B as a single governed turn of the architecture cycle. A scientific subtlety is load-bearing here: the intermediate step targets a local binding site of roughly 24 active electrons (48 qubits) as its scientific ceiling — the precise scale at which both exact classical diagonalisation ($\sim 18e$, where the determinant count turns combinatorial) and full state-vector simulation (~ 4.5 PB) collapse. Operationally the run executes at $\sim 20e/40q$ (~ 17.6 TB, within multi-node HPC reach on NSF ACCESS / USC CARC). In either case this is not classical chemistry: it is the variational quantum algorithm executed on a classical HPC simulator of the quantum circuit. The architectural objective — demonstrating that the governed pipeline scales to meaningful quantum-circuit width — is fully achieved at 40 qubits (see §8, Table 8.1 for the memory-scaling rationale). The provenance layer (P1–P9) operates throughout as TOGAF's Architecture Governance, rendering each iteration auditable. The contribution of this dissertation is located precisely at the intersection of three

otherwise-mature fields — classical HPC, quantum computation, and computational oncology — where no prior system has operated under a unified, provenance-complete architecture.



[Figure — rendered from source]

§ 05 Research Methodology — Five Phases.

The methodology is structured as five phases of Design Science Research (Hevner et al., 2004) — building an artifact, evaluating it against defined criteria, and refining through expert review. The Phase 3 proof-of-concept anchors on superconducting hardware for the empirical-precedent reasons defended in §1.4; the architecture itself remains hardware-agnostic at the workflow level, with vendor-specific compilation and error-mitigation layers documented explicitly (§1.4).

Foundation

Patient NGS analysis · Literature verification · BQP-relevant classification argument · Active space determination (~44e, ~88 qubits)

Platform Prototype

Six-screen SOLANGE orchestration platform · Hardware-agnostic backend abstraction · FDA 21 CFR Part 11 compliance module · 3D molecular visualization · Deployed at solange-platform.bio

Architecture Validation Run · 4-qubit Simulator + HPC Surrogate

Objective: demonstrate that the governed orchestration pipeline executes correctly across simulator and HPC compute tiers before first quantum hardware contact.

Simulator (complete): 347 runs across 15+ mutation targets on PennyLane default.qubit v0.38.0 — including TP53_C275F (116 runs), KEAP1_LOF (39), STK11_LKB1 (44), and twelve additional targets. 4-qubit VQE, CAS(2e,2o)/STO-3G, AllSinglesDoubles UCCSD ansatz, 80 Adam steps — convergence to within 1×10^{-5} Ha of CASSCF-exact, ~7 s on Render cloud. Full P1–P9 audit chain, SHA-256 P8 seal, and Supabase persistence verified 100% PASS across all 347 runs.

HPC (complete): 4 runs on NVIDIA L40S / Laguna (CPU sparse and lightning.gpu backends) — surrogate and native molecules at 8–12 active electrons, validating that the orchestration pipeline, provenance chain, and DMRG classifier remain intact when execution shifts from cloud to GPU cluster. DMRG class=C confirmed ($S_{\max} < 0.5$, classically sufficient at surrogate scale) on all HPC runs. Full audit trail, SHA-256 P8 seal, and Supabase persistence 100% PASS.

DMRG classification (complete): 15 runs across four targets — TP53_C275F (6), ARID2_LOF (5), STK11_LKB1 (2), KEAP1_LOF (2) — establishing each target's classical-tractability class (A/B/C) from measured entanglement ($S_{\max}$) rather than system size, via block2's DMRGDriver on Laguna. CPU-only: this pipeline's DMRGDriver is constructed with a thread count, not a GPU/CUDA argument, so — unlike the HPC row above — no NVIDIA hardware is exercised here even when the node carries one. Full audit trail, SHA-256 seal, and Supabase persistence verified.

Architecture Scaling Run · Real QPU Pipeline + Full Binding Site (~44e/88q)

The objective of Phase 3B is to demonstrate that the architecture scales without re-derivation — first across the classical → quantum hardware boundary, then across the active-space boundary to the full clinical target.

Dry-run (complete): 3 local statevector trials (3B-QPU-dryrun) completed before first real hardware contact — verifying circuit compilation, telemetry, and audit chain behaviour in QPU-mode without consuming QPU credits.

Real QPU (complete — fixed-state measurements): 16 runs on `ibm_marrakesh`, `ibm_kingston`, and `ibm_fez` (Qiskit Runtime EstimatorV2 primitive, real superconducting QPUs) — 15 at 4 qubits (4,096 shots, ~21 QPU seconds each) and one at 8 qubits (CAS(4,4), TP53_C275F mutant, 4,096 shots, 125 QPU seconds; measured $\langle H \rangle = -3.112973$ Ha vs. an exact HF reference of -3.118615 Ha — 5.64 mHa hardware noise, expected at this scale, not an error). Targets: TP53_C275F mutant/native (8 runs), KEAP1_LOF mutant (4),

STK11_LKB1 native (1), CDKN2A_LOF native (1), ARID2_LOF native (1), H₂ smoketest (1). All 16 runs 100% PASS, seal_ok=true, notarized via QPU/IBM ingestion pipeline. Architecture, audit trail, and compliance record verified intact across the classical → quantum boundary — now demonstrated at twice the circuit width.

First on-device optimization (complete — 4-qubit, single parameter): 4 real-hardware jobs on ibm_marrakesh (3-point exact sinusoid fit + 1 confirmation measurement), targeting TP53_C275F mutant with a single-parameter double-excitation ansatz — the first Phase 3B run to search a parameter ON the QPU itself, not just measure a fixed state. Result: confirmation energy -266.46418 Ha, 26.6 mHa from CASSCF-exact — worse than the plain Hartree-Fock reference (18.0 mHa from exact), because with only 3 unaveraged measurements per fit, hardware shot noise dominated the small sinusoidal signal being resolved. Reported as an honest negative result, not hidden: it demonstrates the on-device optimization pipeline runs end-to-end and notarizes correctly, while showing that this minimal-cost fit strategy needs more shots or measurement points before its result is trustworthy at this problem's energy scale.

Full binding site — IBM Heron r3 (~88q): Pending. Re-execute the same orchestration workflow on the full binding site, incorporating the Zn²⁺ coordination environment (C176, H179, C238, C242) and the DNA-guanine interface (~44e/~88 qubits under Jordan-Wigner encoding). The IT-side test is whether the audit trail and compliance record retain integrity across the active-space boundary change. Contingent on IBM Quantum Network institutional access.

The Transitional Window: Situating SqDRIFT and Heron r3 Within IBM's Fault-Tolerant Roadmap

The selection of IBM Heron r3 as the quantum hardware substrate and SqDRIFT as the molecular simulation algorithm reflects a deliberate positioning within a well-defined and time-bounded technological window. IBM's publicly disclosed roadmap — supported by a \$10+ billion capital commitment and a \$1 billion CHIPS Act award — establishes a clear architectural trajectory away from noisy intermediate-scale quantum (NISQ) devices toward fault-tolerant systems based on quantum low-density parity check (qLDPC) error correction. The milestones on this trajectory are specific: the Kookaburra processor (2026) introduces modular quantum memory, Cockatoo (2027) entangles multi-chip configurations, and Starling (2029) delivers the first large-scale fault-tolerant system at 200 logical qubits executing 100 million quantum gates.

This roadmap does not diminish the contribution of the present work — it contextualizes and justifies it. Heron r3, as the current-generation superconducting processor, and SqDRIFT, as a leading NISQ-era sample-based quantum diagonalization protocol whose stochastic state sampling inherits qDRIFT's formally bounded error scaling, represent the state of the art during precisely the window in which clinically relevant quantum molecular simulation first becomes computationally tractable. IBM's own 2026 demonstration (IBM Research; arXiv:2603.08696) applied SqDRIFT on 100 qubits to a 32-electron active space of C₁₃Cl₂ — a benchmark that for the first time brings the active-space requirements of NSCLC tumor-suppressor mutations within reach of available hardware.

IBM's own strategic moves reinforce this trajectory. In July 2026, IBM announced the acquisition of HRL Laboratories — formerly a joint research entity of Boeing and General Motors — establishing a formal two-track qubit roadmap: the existing superconducting qubit line (Heron r3 → Kookaburra → Starling → Blue Jay), which SOLANGE Phase 3B directly targets, and a new silicon-spin qubit track (HRL expertise, paired with the Anderon quantum wafer foundry established by IBM earlier in 2026 with U.S. Department of Commerce support) aimed at the fault-tolerant regime beyond the current NISQ ceiling. SOLANGE's Phase 3B target — C275F at 88 qubits, within the 94-qubit demonstrated ceiling (Merz et al. 2026) — sits squarely on the superconducting track. The KEAP1 and STK11 full-site targets, requiring 300+ qubits, fall in the fault-tolerant window now formally assigned to IBM's silicon-spin trajectory (~2030+). IBM's own acquisition is, in TOGAF terms, an ADM-class architectural migration — and constitutes independent industry validation that the NISQ-to-fault-tolerant transition this dissertation anticipates is a recognised engineering trajectory, not a speculative one [60] [61] .

The significance of this window extends beyond the chemistry. The transition to fault-tolerant quantum computing will not resolve — and in fact will intensify — the information systems challenges that constitute the primary contribution of this dissertation. As IBM scales toward Starling (2029) and the Blue Jay quantum-centric supercomputer (2,000 logical qubits, 2033), targeting drug discovery as an explicit commercial vertical, the regulatory reconstruction problem becomes more acute, not less. Probabilistic, learned quantum outputs will be generated at greater scale and complexity, yet no established framework exists for governing their audit trails, ensuring reproducibility, or satisfying FDA 21 CFR Part 11 requirements in a clinical submission context. The pipeline architecture, P1–P9 provenance schema, and bifurcation-aware orchestration developed in this dissertation are designed precisely for this trajectory: they establish the IS governance substrate that fault-tolerant quantum drug discovery workflows will require, built and validated during the window in which the underlying quantum computation first becomes feasible.

The present work therefore occupies a strategically necessary position: it is neither premature nor transitional. It is foundational.

Triangulated DSR Evaluation

Four converging evidence sources (technical, comparative, expert, case-based), executed in parallel sub-phases 4A–4F over Jan–Apr 2027. Detailed below.

Writing & Defense

Dissertation drafting May–Sep 2027 · Pre-defense committee revisions Oct 2027 · Target defense Q4 2027

Phase 4 · Triangulated DSR Evaluation Architecture

The artifact is evaluated through four independent evidence sources, following the FEDS framework (Venable et al., 2016) for DSR evaluation strategy. Each layer addresses a different validity threat; together they triangulate to produce a defensible aggregate judgment that no single source could supply alone.

Layer 1 · Technical Evaluation (Artifact Performance — No Human Subjects)

Objective measurement of the artifact itself, executed via an automated test harness with pre-registered acceptance thresholds.

Criterion	Metric	Acceptance Threshold
End-to-end workflow execution	% of pipeline stages completing without manual intervention across 50 simulated runs	≥ 95%
Audit-trail integrity (P1–P9)	% of runs producing complete, cryptographically valid provenance records	100%
Cross-vendor reproducibility	Δ between PennyLane and Qiskit Aer simulator outputs (χ^2 test on outcome distributions)	$p > 0.05$
Compilation-layer translation overhead	SWAP-gate count delta between superconducting and trapped-ion compilation targets	Documented & bounded
Active-space scaling integrity	Audit-trail completeness across local (~24e) → full (~44e) boundary change	100%
Compliance record validation	% of P1–P9 records passing automated Part 11 §11.10 validation checks	100%
Time-to-result	Wall-clock from NGS input to validated output, simulator backend	Measured & reported
Failure-mode coverage	% of injected fault scenarios (network drop, calibration drift, sampling truncation) detected by audit layer	≥ 90%

Layer 2 · Comparative Evaluation (Benchmarking vs. Existing Alternatives)

Reference task — a VQE simulation of a small molecule (H_2 or LiH) — executed through this platform and through three baseline pathways: Qiskit Runtime alone, Orquestra (Zapata), and fully manual workflow. Comparison metrics: number of manual steps, setup time, and whether the output satisfies 21 CFR Part 11 audit requirements.

Capability	This Platform	Qiskit Runtime	Orquestra	CUDA-Q + NVQLink	Manual
Multi-vendor QC backend abstraction	✓	✗ (IBM only)	Partial	✓ (cross-vendor, excludes IBM)	✗
HPC integration (GROMACS / ORCA)	✓	✗	✗	✓	Manual
AI rescoring layer	✓	✗	✗	Partial (CUDA-Q-native AI primitives)	Manual
21 CFR Part 11 audit trail	✓ (P1–P9)	✗	✗	✗	✗
Stakeholder visualization layer	✓	✗	✗	✗	✗
Statistical-equivalence validation	✓	✗	✗	✗	✗
Bifurcation-aware (across IBM-native & Nvidia substrate)	✓	✗	✗	✗ (one side of the bifurcation)	✗

Layer 3 · Expert Panel Evaluation (Stakeholder Validity)

Panel of $n = 4\text{--}6$ experts, composition deliberately calibrated for methodological coverage:

- 1 quantum chemistry / quantum computing expert (academia) — scientific and algorithmic rigor
- 1 oncology / molecular pathology clinician — clinical relevance and target validity
- 1 pharma R&D / computational drug discovery (industry) — translational practicality
- 1 Information Systems / DSR methodologist — methodological-validity guard (Hevner G5)
- 1–2 additional reviewers (flex) — FDA / regulatory affairs or hospital IT architect / clinical informatics

Composition aligned with dissertation-committee guidance ($n \leq 6$) and mapped to the Triangulated DSR layers: each layer of evaluation has at least one independent reviewer, with the methodologist guarding rigor across all four.

Evaluation protocol follows Hevner et al. (2004) — the Seven Guidelines, mapped to concrete instruments:

Guideline	Evaluation Question	Method
G1 · Design as artifact	Does the artifact exist, documented and operational?	Live demo + code review
G2 · Problem relevance	Is the problem significant to stakeholders?	Likert survey + open-ended response
G3 · Design evaluation	Does the artifact meet utility / quality / efficacy?	Triangulated (Layers 1–3)
G4 · Research contributions	Are the contributions novel and defensible?	Peer review by domain experts
G5 · Research rigor	Is the methodology rigorous?	DSR methodologist review
G6 · Design as search process	Are rejected alternatives documented?	Design-log review
G7 · Communication	Is the artifact transferable to stakeholders?	Stakeholder usability sessions

Scoring: Likert 1–7 on each criterion. Acceptance threshold: mean ≥ 5.0 with no individual score < 4 . Inter-rater reliability: Krippendorff's α reported on agreement-targeted criteria; $\alpha < 0.67$ flags a criterion requiring refinement or signaling genuine disagreement.

Layer 4 · Case-Based Walkthrough (Peffer et al., 2007)

Two clinical scenarios are traced end-to-end through the platform, one real (the anchoring patient — MI25-0349 with TP53 C275F), one hypothetical (a KEAP1 loss-of-function case). For each scenario the documentation captures: every design decision invoked; every point at which the audit trail caught — or should have caught — an anomaly; every step requiring manual intervention; and the stakeholder output produced for each persona (oncologist view, pharma view, IT view). This narrative complements the numbers — the case-based evidence is what shows the system works in context, not only on the bench.

Phase 4 Sub-Phases & Schedule

Sub-phase	Activity	Duration
4A	Layer 1 — automated technical-evaluation harness	4 weeks
4B	Layer 2 — comparative benchmarking	3 weeks
4C	Expert-panel recruitment & briefing materials (parallel to 4A/4B)	3 weeks
4D	Panel sessions — asynchronous, two rounds (theoretical + hands-on)	4 weeks
4E	Case-based walkthrough documentation	2 weeks
4F	Triangulation analysis & evaluation report	3 weeks

Total ~14–16 weeks, fitting the Jan–Apr 2027 window with 4A/4B/4C running in parallel.

Pre-Registration of Evaluation Criteria

All Layer-1 acceptance thresholds, Layer-3 Likert criteria, and Layer-4 walkthrough scenarios are pre-registered as a frozen document on OSF (Open Science Framework) prior to the start of Phase 4. This neutralizes the post-hoc-fit critique — that thresholds were adjusted after results were observed — and aligns the dissertation with the FEDS evaluation framework (Venable, Pries-Heje & Baskerville, 2016).

External Validity Statement

The evaluation is run on one mutation instance (TP53 C275F) with one quantum modality anchor (superconducting). Generalization to other mutations or other modalities is therefore analytical generalization in Yin's sense (2014) — an inference about the scope of the artifact's design principles, not statistical extrapolation from a sample. The architecture's mutation-agnostic and (workflow-level) hardware-agnostic design supports analytical generalization to BQP-relevant fermionic targets and gate-model quantum backends; statistical generalization is explicitly outside the dissertation's scope.

§ 06 Original Contributions — Three Claims.

This chapter advances three original Information Systems contributions. The biochemical setting — TP53 C275F, non-druggable tumor-suppressor mutations, electronic structure of binding sites — provides the anchor problem class, but the contributions themselves are IS-research artifacts: a classification framework, an orchestration architecture, and a regulatory-compliance schema. The dissertation does not aim to solve the biochemical problem of drug discovery directly; it builds the systems-engineering scaffold within which such work becomes governed, reproducible, and reusable.

Novelty-verification note. Each "first"/"no published work" claim in this chapter is stated as of June 2026. Priority was checked against Google Scholar, arXiv (quant-ph, cs.CE), PubMed, IEEE Xplore, and the Web of Science Core Collection, using combinations of the terms quantum · 21 CFR Part 11 / provenance / audit trail · drug discovery / oncology / tumor suppressor · orchestration / QC-AI-HPC, with a literature cutoff of June 2026. No prior work combining all elements of a given claim was found at that date; each claim is bounded to that evidence base and is revisable as the literature advances.

BQP-Relevant Classification Framework for Non-druggable Tumor Suppressor Targets

The first published analysis (as of June 2026) establishing TP53 C275F as a BQP-relevant problem in the context of oncology drug discovery — under the standard complexity-theoretic framing of quantum chemistry (Abrams & Lloyd, 1999; Aspuru-Guzik et al., 2005). The framework defines the active-space boundary (~44 electrons) beyond which the best known classical methods do not achieve the required chemical accuracy within tractable resources, and at which a known BQP-class quantum algorithm provides polynomial-time tractability. The framework is dual-conditional: classical intractability in practice plus BQP-class quantum tractability in principle. It does not claim BQP-hardness or a formal classical-vs-quantum separation, both of which remain open in general; it claims a defensible operational regime under which quantum-derived molecular evidence is the appropriate target of an orchestration architecture. The framework is explicitly conditional on the current state of classical methods (DMRG, neural quantum states, tensor networks) and revisable as those methods advance — providing a reusable classification template applicable to other non-druggable mutation targets. Operationally, the template sorts targets into three classes — the same classification the SOLANGE system applies in its live Non-Druggable Target Registry, where every mutation ingested from an NGS report is mapped and prioritized on arrival. The rule reads two measured quantities: the number of active electrons at the target's site, and the maximum bipartite entanglement entropy $S_{\max}$ measured over that site. At or below 18 active electrons the target is Class C whatever its entanglement, and entanglement is not consulted at all — exact diagonalisation builds the full configuration space and its cost depends only on how many configurations there are, so below that wall no degree of entanglement makes a target hard. Above 18 active electrons the entanglement decides: $S_{\max}$ at or below the threshold is Class B (quantum-advantaged), where classical tensor-network methods still converge and quantum simulation offers an advantage rather than a necessity, and $S_{\max}$ above it is Class A (quantum-necessary), where the bond dimension required outruns what any classical method delivers. The classification is thus a routing and prioritization policy — Class A and B enter the quantum queue; Class C is filtered out.

A class follows only from measurement. Electron count is known for every target from its structure; $S_{\max}$ requires a run covering at least 90% of the active site, the Gateway's scope gate (§06.ii). Where that run does not exist, no class is asserted: the entry carries the expected class as a prior, together with a note naming the measurement that would settle it. A prior is a hypothesis on the record, not a result. Electron count may motivate a prior; it cannot discharge one. Four STK11 entries are so marked, their Class A resting on the size of the site rather than on any entanglement measurement, and C275F's Class B is likewise a prior until its site is measured at full coverage.

Note — why two quantities rather than one. The cost of the tensor-network methods that define the classical frontier factorises into two independent terms: an exponential in the entanglement, $e^{3 \cdot S}$, which contains no system size at all, multiplied by a polynomial in the

orbital count, which contains no entanglement at all. Neither term settles tractability alone. A small but strongly entangled system stays cheap because the polynomial term is small — stretched N_2 carries $S_{\max} = 2.86$, far above the threshold, and solves in seconds. A large but weakly entangled one can be costly and still classically solvable, as DMRG at CAS(113,76) demonstrated for FeMoco. Only when both terms are large does the classical path close, which is why the classification reads both axes. The 18-electron boundary is the point past which exact diagonalisation becomes combinatorially impractical; the entanglement threshold is calibrated against N_2 at equilibrium ($S_{\max} = 0.42$) and stretched ($S_{\max} = 2.86$), and is a policy of this pipeline rather than a constant of nature — it marks where this pipeline’s committed bond dimension runs out, not where the method itself stops. Under this template C275F is Class B: it sits in the empirically quantum-advantaged regime defined by the current state of both classical and quantum methods circa 2026. This is a theoretical IT contribution; it does not require hardware execution to be valid.

The physical basis of the classification — entanglement, not electron count. The ~ 44 -electron figure names the size of the active space the anchor target requires; it does not by itself explain why that space is hard, and it is not the quantity the framework classifies on. What governs classical tractability is the entanglement structure of the wavefunction over the active space. In the tensor-network methods that define the current classical frontier, the wavefunction is carried as a matrix product state whose retained rank — the bond dimension M — is bounded above by e^S , where S is the von Neumann entanglement entropy across the most strongly correlated bipartition of that space (written $S_{\max}$), while cost grows as M^3 . That bound follows from assuming a flat entanglement spectrum; real spectra decay, so the requirement rises more slowly than the bound. Calibrating it directly on an N_2 stretch series at CAS(10,16) — recording $S_{\max}$ and the bond dimension actually required for chemical accuracy at each geometry — returns a growth exponent of 0.31 rather than 1 (five points, $R^2 = 0.82$). The dependence remains exponential in entanglement and free of system size, which is what the classification rests on; its rate is milder than the bound. A target whose $S_{\max}$ saturates as the active space is enlarged toward full site coverage — an area law — therefore remains classically tractable however many electrons it nominally contains; one whose $S_{\max}$ keeps climbing — a volume law — does not, however much classical hardware is applied to it. The published DMRG treatment of FeMoco at CAS(113,76) and bond dimension 6000 (Brabec et al., 2021) settles the point empirically: a 113-electron active space proved classically solvable because its entanglement saturated, and that it saturated was established by performing the calculation, not predicted from its size. The corresponding failure of single-reference methods such as CCSD(T) is likewise a failure under strong static correlation rather than at a particular electron count; the two coincide often enough across main-group organic chemistry for size to serve as a working proxy, and they part company precisely in the multireference regimes that non-druggable tumour-suppressor sites occupy. Electron count is retained throughout this dissertation as a resource-planning variable — it fixes qubit width, memory footprint, and scheduling — but the classification criterion is measured entanglement.

This is not a framing applied to the results after the fact; it is the rule the classifier executes. The DMRG rung reads $S_{\max}$ directly off the converged matrix product state and classifies against a calibrated threshold ($S_{\text{HARD}} = 1.5$, sited between N_2 at equilibrium, $S_{\max} = 0.42$, and stretched N_2 , $S_{\max} = 2.86$), together with whether DMRG reached chemical accuracy within a practical bond dimension. A target below the exact-classical wall is returned Class C on the ground that exact diagonalisation already suffices; above that wall, a target that converges at practical M with low entanglement falls in Class B, and one whose $S_{\max}$ exceeds the threshold — so that the bond dimension demanded at the full site grows beyond reach — falls in Class A. Electron count enters only in the first of those three tests, and there only as the boundary of exact diagonalisation, whose cost genuinely is combinatorial in the size of the space. The Computational Decision Gateway (§06.ii) accordingly accepts a measured entanglement value as evidence and declines to classify on size alone, and its scope gate requires a measurement to cover at least 90% of a target’s active site before it counts at all — an entropy measured over a fraction of a site says nothing about whether $S_{\max}$ saturates across the whole of it. No full-coverage measurement yet exists for C275F. Its Class B designation

therefore records the present state of the evidence, and is revisable in either direction by a measurement that does cover the site: an $S_{\max}$ that saturates would argue it toward the classically sufficient regime, one that keeps climbing would argue it to quantum-necessary.

The framework is designed as a reusable classification tool applicable beyond the four anchor mutations. Fifteen pipeline-eligible non-druggable mutations have been identified across eleven genes (TP53, KEAP1, CDKN2A, STK11, NF1, RB1, SMARCA4, RBM10, ARID1A, SETD2, U2AF1), collectively affecting an estimated 1.1–1.3 million unique NSCLC patients annually — representing the largest unaddressed population in computational oncology within the quantum feasibility window. Each satisfies the dual BQP-relevant condition: classical intractability at chemical accuracy at the specified active-space boundary, and a known BQP-class quantum algorithmic formulation for the corresponding fermionic Hamiltonian eigenvalue problem. The classification template defined here for C275F is directly applicable to all fifteen.

Part-11-Governed QC·AI·HPC Orchestration for Oncology Mutation Simulation

A cross-vendor orchestration architecture designed specifically for oncology tumor-suppressor mutation simulation, in which 21 CFR Part 11 provenance is emitted at workflow execution time rather than reconstructed retrospectively. As established in §3.4, the post-April-2026 landscape contains a generic cross-vendor substrate (CUDA-Q + NVQLink) that this architecture sits on top of where IBM's path permits, and across where it does not — Qiskit Runtime when running on IBM Heron, CUDA-Q-equivalent paths if a non-IBM modality is later substituted (§1.4.1). The contribution is not the cross-vendor abstraction itself, which is now commodity, but three properties the commodity substrate does not provide:

(1) Oncology-mutation-specific construction. The pipeline is built around NGS-driven active-space selection, mutation-aware Hamiltonian assembly, and AI rescoring of binding outputs for tumor-suppressor targets — not generic small-molecule chemistry of the kind exercised by IonQ/AstraZeneca on Suzuki–Miyaura coupling or Quantinuum/Nvidia on imipramine.

(2) Workflow-embedded Part-11 governance. The eight-element provenance record specified in §06.iii (extended to a ninth element P9 covering ML-decoded QEC, per §3.5) is emitted as a first-class artifact of workflow execution — capturing calibration epoch, per-gate error budget, and raw outcome distribution at the moment they are recoverable. Adding such records retrospectively is non-physical for probabilistic quantum simulation; the governance must be in the orchestration, not bolted onto its outputs.

(3) Selective paradigm invocation, bifurcation-aware. The workflow abstraction encodes a decision rule for when subproblems are dispatched to AI, HPC, or QC backends — based on problem characteristics such as active-space size, classical-method tractability under §06.i's BQP-relevance classification, and ML-decoder availability. The same abstraction does not commit the dissertation to a single vendor stack as the IBM Quantum-Centric Supercomputing path and the Nvidia NVQLink/CUDA-Q substrate continue to diverge; re-anchoring on a different modality, should the empirical landscape shift, requires re-implementing only the compilation and error-mitigation layer — not the workflow. The orchestration is therefore not merely cross-paradigm execution but selective computational escalation : AI handles pattern-detection and rescoring; HPC handles coarse classical simulation and conformational filtering; QC is invoked only where classical methods exceed practical limits at the active-space boundary. The platform's operational interface makes this escalation legible rather than implicit: the Orchestration surface is laid out as a top-to-bottom ladder whose rungs climb in problem size — an in-browser proof run on a laptop (Rung 1), classical CASSCF/VQE dispatched to the HPC cluster (Rung 2), the DMRG tensor-network classifier that establishes the last classical resort (Rung 3), and real quantum hardware at the foot of the ladder (Rung 4). Each rung carries its own internal sub-ladder, and a target descends only as far as its active-space size and BQP-relevance class require, so the reader of the running system sees the same decision rule the workflow encodes.

The orchestration remains hardware-agnostic at the workflow level and vendor-specific at the compilation and error-mitigation layer , documenting — rather than concealing — the translation overhead these layers require. The first cross-vendor, Part-11-governed

orchestration architecture (as of June 2026) designed specifically for oncology mutation targets. Validated by proof-of-concept execution on PennyLane simulator. Designed for IBM Heron r3 execution in the subsequent research phase.

A further consequence of this architecture is frequently underweighted in comparisons against precedent systems: the absence of inter-tool hand-offs is itself a governance property, not merely an operational convenience. The workflow described above — NGS-driven active-space selection, mutation-aware Hamiltonian assembly, classical reference computation, quantum circuit execution, error mitigation, and AI rescoring — would, absent a unifying orchestration layer, require a researcher to coordinate five or six independently maintained tools (a classical electronic-structure package, a quantum SDK specific to the target vendor, a separate quantum-sampling toolchain for Phase 3B, an HPC scheduler, and an ad hoc record-keeping system reconstructed after the fact). Each hand-off between such tools is a point at which the provenance chain specified in §06.iii can break silently: a result re-entered by hand, a parameter copied without its calibration context, a run whose device state was never captured because the capturing tool did not exist at that boundary. The single-platform design is therefore not offered as a usability feature distinct from the Part-11 claim of §06.iii — it is the precondition for that claim. P1–P9 records can be emitted as first-class artifacts of workflow execution (§06.ii) only because no boundary exists at which provenance capture is optional or manual. None of the precedent systems reviewed in §3.4 — including the architecturally proximate Cleveland Clinic/IBM QCSC workflow — operate as a single governed platform across this full pipeline; each composes independently maintained tools without a provenance-preserving substrate spanning the composition.

Secure cross-boundary dispatch (pull-based orchestration). A governed HPC environment such as USC CARC's Laguna cluster enforces interactive multi-factor authentication (Duo) at every login and prohibits automated inbound connections from external services — a security posture typical of institutional research computing. This forecloses the intuitive design in which the orchestration platform pushes jobs into the cluster: no credential an external service could store would satisfy an interactive 2FA challenge, and storing one would itself violate the cluster's access policy. The architecture therefore inverts the direction of control. Rather than the platform reaching into the secured environment, a bounded execution agent running inside the researcher's already-authenticated cluster session polls the platform's dispatch queue over an outbound HTTPS channel — the only channel the security boundary permits — claims a queued job, executes it with full in-environment privileges, and returns the verified P1–P9 provenance record through that same outbound path (§06.iii). This is the established self-hosted-runner and message-queue-worker pattern from continuous-integration and distributed-systems practice, adapted to governed scientific computing: it keeps zero cluster credentials outside the trust boundary, opens no inbound path, and does not circumvent the human 2FA that gates the session. Crucially, the dispatch itself becomes a provenance-tracked lifecycle — queued → claimed → executed → returned — so that even the act of routing work across the QC·AI·HPC fabric is auditable rather than opaque. That visibility extends into execution itself: while a claimed job runs, the agent reports its current stage back to the same dispatch row over the outbound channel — the QPU's live IBM-side state (queued → running on the device), and on the classical rung the progression through RHF → CASSCF → VQE optimisation step or, for DMRG, successive bond dimensions — so that a long-running job surfaces live progress in the interface rather than appearing as an opaque “running” with no indication of where it is. This is the same audit-in-running-code discipline applied to the in-flight state, not only the returned record. The honest limitation is symmetric to the security guarantee: work executes only while an agent is active within a session; absent one, dispatched jobs remain queued by design rather than executing through some privileged back channel. This limitation is surfaced rather than hidden in the operational interface: the dispatch panel shows a live agent-liveness indicator (green when an agent is heartbeating, red when none is), and when it is red the platform offers a “wake agent” affordance that hands the researcher the exact launch command to paste inside their own authenticated session — it deliberately does not start the agent remotely, because doing so would require the very inbound channel the architecture forecloses. The button is therefore a usability

convenience that reaffirms the boundary rather than eroding it: even the recovery path from an offline agent runs through the human-authenticated session, never around it.

The boundary protects both parties' intellectual property, not only the customer's data. DP2's pull architecture is presented above as a data-security guarantee — no cluster credential, and by construction no patient record, sequence, or structure, ever needs to leave the customer's own perimeter, because the agent that touches that data runs inside it. The same boundary has a second, symmetric direction that a commercial deployment must make explicit: the Computational Decision Gateway (§06.i) that recommends a compute route is deliberately scoped to accept only scalar evidence — an active-electron count, a measured entanglement value, a correlation energy, an opaque target label — never a structure, a sequence, or a record. A customer's proprietary molecule therefore never has to cross into the Gateway to receive a routing recommendation, and the Gateway's own scoring function — the thresholds, the provenance-ceiling logic, the rationale generator — never has to cross out to the customer to produce one. What is exchanged is a route, a confidence, and a written rationale; what is exchanged is not the customer's molecule and not the platform's decision logic. This is the condition under which a hybrid on-prem deployment (§08) is commercially viable at all: a customer who cannot allow proprietary compounds to leave their firewall can still consult the Gateway, and the party operating the Gateway can still protect the scoring logic as a trade secret rather than disclosing source. The property is not a separate mechanism bolted on for this purpose — it is the same DP2 boundary already defended above, read in the direction that matters to a prospective licensee rather than to a security reviewer.

The notarization boundary — LEON. A returned result crossing back into the platform is, at that instant, an unverified external claim : the dispatch inverted control precisely so that no credential and no trust would live outside the cluster's boundary, and that same logic must apply to the payload the agent returns. The architecture therefore terminates the return path at a dedicated notarization component — LEON (Lineage-Evidence Orchestration & Notarization). LEON does not accept a submitted result on the strength of its origin; it independently recomputes the P8 cryptographic seal over the P1–P7,P9 fields, compares it against the seal the agent transmitted, re-checks the internal physics-consistency invariant (core-plus-active energy against the reference), and rejects any record that fails — the row never enters the store. The same notary re-runs on demand at query time, so a record's integrity is re-attestable long after ingestion rather than asserted once and trusted thereafter. This operationalizes a single design principle — verify, don't trust : within a governed pipeline, trust in a result is derived from reproducible cryptographic proof, never from the identity of the boundary it arrived through. LEON is the component that makes the provenance chain of §06.iii self-enforcing rather than declarative, and it is the element of this architecture with no counterpart in the precedent systems reviewed in §3.4 — an independently re-verifiable notarization guard sitting on a cross-security-boundary return path under a Part-11 mapping.

The P8 seal, concretely. §06.iii defines the P1–P9 schema in the abstract; the table below is its literal, field-level instantiation in the running Phase 3A/3B implementation — the exact set LEON's SHA-256 computation runs over. Every field whose name is prefixed p1_ through p7_ , plus p9_ , is serialized into one canonical JSON object (keys sorted, so the same data always produces byte-identical input regardless of field insertion order) and hashed. p8_* fields are the seal's own output and are necessarily excluded from what it hashes over — a seal cannot certify itself without becoming circular. One further field, p3_calibration_epoch , is deliberately excluded despite its p3_ prefix: timestamp values are reformatted in a lossy way on the database round-trip, so including it would make a query-time re-verification fail on formatting alone rather than on an actual integrity violation — the exclusion keeps the seal robust to storage-layer normalization while still covering every field that carries scientific content.

Segment	Representative fields (Phase 3B/QPU record)	What it captures
P1	p1_qubit_count , p1_ansatz , p1_circuit_hash	Circuit identity — what was actually run
P2	p2_encoding , p2_basis_set , p2_active_electrons , p2_active_orbitals , p2_model_compound , p2_jw_terms	Problem definition — the full Jordan–Wigner Pauli-term list, not a summary
P3	p3_backend , p3_backend_version	Hardware identity (p3_calibration_epoch excluded — see above)
P4	p4_gate_error_rate , p4_readout_error_rate , p4_t1_us , p4_t2_us	Live device-noise telemetry at execution time
P5	p5_shots , p5_elapsed_s , p5_ecore_ha , p5_active_energy_ha , p5_casscf_ref_ha	Computation internals — shot count and intermediate energetics
P6	p6_method , p6_note	Method/algorithm identity
P7	p7_energy_ha , p7_ref_hf_ha , p7_ci_lower , p7_ci_upper , p7_confidence	Reported result — point estimate plus its 95% confidence interval, not a bare number
P8	p8_hash , p8_seal_payload , p8_algorithm	The seal itself — output, not input, of the hash
P9	p9_applicable , p9_note	ML-decoder lineage (forward-looking; deterministic-decoding default in Phase 3)

Field names shown are from the QPU/Phase 3B record shape (`solange_qpu.py`); the classical Phase 3A/HPC record shares the same P1–P9 prefix convention with a slightly different field set, sealed by the identical mechanism in `backend/routes/leon.py` .

FDA 21 CFR Part 11 Compliance Framework for Probabilistic Quantum Simulation

21 CFR Part 11 was drafted under the assumption that computational outputs are deterministic and reproducible bit-for-bit. Quantum simulation breaks that assumption at a foundational level: identical input circuits, executed on the same physical device under nominally identical conditions, yield distributions of measurement outcomes, and those distributions themselves shift with hardware noise, decoherence, and calibration drift. The April 2026 emergence of ML-decoded quantum error correction (Nvidia Ising) introduces a second-order break: the decoder itself is now a learned, probabilistic model whose lineage must enter the compliance record (§3.5). This contribution closes the regulatory gap with a defensible specification of what records must be retained, in what form, for quantum simulation outputs to satisfy Part 11.

Five Sources of Output Variability

The framework enumerates the irreducible sources of variance in a quantum simulation output, each with distinct provenance characteristics: (i) statistical sampling variance ($1/\sqrt{N}$ convergence of estimators over shot count — a property to be documented, not eliminated); (ii) hardware noise (gate infidelity, readout error, crosstalk — captured at execution time, not retrospectively); (iii) decoherence (T_1/T_2 recorded as device-state metadata at the calibration epoch); (iv) calibration drift (the boundary condition that defines, in operational terms, which physical device executed the simulation); (v) ML-decoder inference variance (when QEC is performed by a learned model rather than a deterministic algorithm — Ising-class decoders introduce inference-time precision effects, batch-ordering sensitivity, and version drift across model retrains, each requiring distinct provenance treatment).

Nine-Element Provenance Record (P1–P9)

Compliance is achieved through a nine-element record attached to every simulation output that may be cited in regulatory submissions. P1–P8 are required for any quantum simulation execution; P9 is invoked whenever the QEC layer involves a learned decoder (e.g., a Nvidia Ising-class pre-decoder or successor models), in which case it records both the learned pre-decoder and the conventional global decoder it escalates to, and degrades to a minimal record (algorithm name + software version) for purely deterministic decoding (pyMatching, BP-OSD). Phase 3 of this dissertation runs on Qiskit with deterministic decoding and does not execute ML-decoded QEC at all; P9 is therefore a forward-looking design contribution — a specification of the records a regulated workflow will require once learned decoders enter production — not a component exercised by the present proof-of-concept.

P1	Circuit fingerprint — cryptographic hash of the post-transpilation compiled circuit
P2	Compilation lineage — source Hamiltonian, fermion-to-qubit mapping, Trotter order/step, SqDRIFT sampling parameters
P3	Device identifier & calibration epoch — backend, processor revision, calibration timestamp
P4	Per-gate error budget snapshot — 1- and 2-qubit fidelities, readout-error matrices, T_1/T_2 at the calibration epoch
P5	Shot count & raw outcome distribution — full bitstring-level histogram, permitting independent re-derivation of any aggregated estimator
P6	Error-mitigation lineage — every transformation applied to raw outcomes (ZNE, PEC, twirling), with parameters and method versions
P7	Statistical estimator & confidence interval — final reported observable, estimator function, 95% CI bounded by shot-count and noise-model variance
P8	Cryptographic seal — time-stamped digital signature over P1–P7 and P9, anchoring the record against subsequent modification
P9	ML-decoder lineage — when QEC inference uses a learned pre-decoder: decoder identifier & version, weights checksum, training-corpus manifest hash, inference-time configuration (precision, batch size, decoder threshold), execution environment (GPU class, CUDA/runtime version), and the conventional global decoder it escalates to (algorithm + version). For purely deterministic decoding: algorithm name and software version. Forward-looking — not exercised in Phase 3.

Why This Differs from Classical Monte Carlo

A natural objection is that classical computation already produces probabilistic outputs: Monte Carlo simulations are stochastic, and Part 11 already accommodates them through pseudo-random-number-generator seed logging — record the seed and the run is reproducible bit-for-bit. What, then, is new here beyond documenting an RNG? The answer is that quantum stochasticity is not seed-controllable. A classical Monte Carlo trajectory is fixed by a single logged integer, so re-execution reproduces it exactly on demand. A quantum run is not: hardware noise, decoherence, and calibration drift are physical processes of the device at execution time, not a controllable parameter, and no logged value can reproduce an identical run even in principle — re-executing the same circuit on the same processor an hour later samples a different noise realization, and after a recalibration it is, for provenance purposes, a different device. This is precisely why §11.10(a)'s bit-for-bit reproducibility test is broken by quantum simulation in a way it is not broken by Monte Carlo, and why the statistical-equivalence reinterpretation below is a genuine extension of Part 11 rather than a restatement of existing seed-logging practice. The ML-decoder layer (P9) is a hybrid case: the learned decoder's weights and version are loggable and reproducible, but its inputs — live syndrome streams from a noisy device — are not; P9 therefore records the reproducible half (model lineage) alongside the irreproducible half (the device state that produced its inputs).

The Critical Reinterpretation

The framework's most consequential move concerns the meaning of "validation" under §11.10(a). For deterministic systems, validation tests verify that re-execution produces identical outputs. For probabilistic quantum simulation this test is non-physical and would mark every legitimate run as anomalous. The framework redefines validation as statistical equivalence: a re-execution is consistent with the original record if and only if the new outcome distribution lies within the confidence interval declared in P7, conditional on the same calibration epoch (P3), circuit fingerprint (P1), and decoder lineage (P9). Where any of these differs — calibration epoch, circuit, or decoder — the comparison becomes a hypothesis test on distribution overlap accounting for declared inter-condition variance drift. This is not a relaxation of the regulation; it is the only operationalization of §11.10(a) consistent with quantum mechanics — preserving Part 11's core intent (detecting unauthorized modification, system failure, or invalid records) while specifying the statistical machinery that must replace bitwise comparison.

Implementation Status — Audit Trail in Running Code

The §11.10(e) requirement for a secure, computer-generated, time-stamped audit trail is not left at the level of specification: it is realized in the deployed prototype (§07) by LEON (§06.ii). Every notarization event LEON performs — a run accepted at ingestion, a stored record re-verified on demand, or a submitted record rejected for a seal or consistency failure — appends one immutable row to a dedicated, append-only audit table (insert-and-read only; no update or delete path exists through its access policies), capturing the

timestamp, the run identifier, the integrity verdict, the recomputed and stored seal hashes, and the actor. The audit trail is therefore generated by the system as a side-effect of verification rather than maintained by hand — the property §11.10(e) is written to guarantee. The same computer-generated date-and-time stamp is surfaced in the operational interface itself, not left in the backing table: every run that returns to the orchestration ladder — HPC (Rung 2), DMRG (Rung 3), or QPU (Rung 4) — carries a fixed timestamp column beside its LEON seal, so the §11.10(e) requirement that the record independently register when each action occurred is inspectable at a glance in the running system. Two honest scope notes apply. First, the audit machinery is exercised in Phase 3A against the classical HPC proxy (§07); its records reduce to the minimal deterministic form, and the P9 ML-decoder lineage remains forward-looking (§08). Second, this contribution addresses the records controls of Part 11 (§11.10(a)–(e)); the electronic-signature controls of Subpart C and the broader regulatory-science questions of IND admissibility are explicitly out of scope (§08). Within that boundary, the audit-trail requirement moves from a design claim to a demonstrated, inspectable artifact.

06.iv · Design Contribution: Principles & Evaluation

The three claims above are, in Design Science Research terms, the instantiation (Hevner et al., 2004; Gregor & Hevner, 2013): SOLANGE is the running artifact that proves they are realizable together. The transferable contribution to Information Systems, however, is not the platform but the smaller set of design principles the platform embodies — knowledge that generalizes to the wider class of governed scientific-computing pipelines that must cross organizational, security, and hardware boundaries while remaining auditable . This subsection states those principles explicitly, positions them against the precedent systems of §3.4, and — in keeping with the honesty commitments of §08 — separates what the proof-of-concept has already demonstrated from what remains for formal evaluation in DSR Phase 4.

Five Design Principles (DP1–DP5)

Each principle is stated as generalizable design knowledge, paired with its concrete mechanism in SOLANGE and the reasoning that justifies it. The principles are deliberately artifact-independent: a different platform in a different scientific domain could satisfy all five without sharing a line of SOLANGE's code.

#	Design principle	Mechanism in SOLANGE (instantiation)	Grounding / rationale
DP1	Verify, don't trust. A result crossing a trust boundary is an unverified claim until reproducible cryptographic proof re-establishes its integrity — never trusted on the strength of its origin.	LEON recomputes the P8 seal and re-checks physics-consistency at ingestion and again on demand; failing records are rejected, not stored (§06.ii).	Applies the zero-trust security stance to scientific provenance; makes the P1–P9 chain self-enforcing rather than declarative.
DP2	Cross a security boundary without weakening it. Where a resource sits behind interactive 2FA with no inbound path, control is inverted so the platform never reaches in and stores no credential outside the boundary.	Pull-based dispatch: a bounded agent inside the authenticated session polls an outbound queue, executes, and returns results — queued → claimed → executed → returned (§06.ii).	Adapts the self-hosted-runner / queue-worker pattern to governed HPC; preserves the human 2FA that gates the session.
DP3	Provenance is a first-class, execution-time artifact. Audit records are emitted at the moment device state is recoverable, never reconstructed afterward.	The nine-element P1–P9 record is produced during execution and sealed at source; calibration epoch and error budget are captured in-band (§06.iii).	Retrospective provenance is non-physical for probabilistic quantum runs; governance must live in the orchestration, not be bolted onto outputs.
DP4	Governance is workload-agnostic. The orchestration and notarization layers are indifferent to the scientific payload they carry.	The CASSCF/VQE chemistry is one interchangeable workload; the dispatch, provenance, and LEON layers make no assumption about it and would carry any heavy compute unchanged.	This is what makes the contribution an Information Systems one rather than a computational-chemistry one — the governed pattern transfers beyond the biology.
DP5	A result's standing is bounded by its weakest input, and the bound travels with the result. A system that computes correctly may still report misleadingly, because correctness and standing are different properties. Every value therefore carries a grade derived from its inputs — never asserted by whoever produced it — and that grade accompanies the value into every surface it reaches.	Evidence grades (demonstration_only / template_conformant / reference_exemplar) are computed from the inputs themselves and cannot be self-declared; the Decision Gateway caps a recommendation's confidence at the provenance of whichever input decided the route, and reports the uncapped value alongside so the gap is visible; a measurement is authoritative only for the active space it actually covered; and the qualification is emitted into the interface, the CSV and JSON exports, and the API payload from one canonical definition, so it cannot be separated from the numbers by a screenshot or a download.	Derived from failures observed in this artifact rather than proposed in the abstract. A correlation energy computed over a two-electron model compound was rendered as a property of a forty-four-electron protein site; because model compounds are shared across genes, four targets displayed the identical value and read as independent corroboration of one another. A DMRG run over a six-electron fragment reclassified the anchor target it had never characterised. Each was a correct calculation reported without its scope. The generalizable claim is that in any pipeline where cheap proxies stand in for expensive systems — which is to say most computational science — scope must be enforced structurally, because it will not survive being left to the reader.

On DP5 and the standing of this platform's own numbers. DP5 is the principle this artifact had the most occasion to learn, and it is stated here with the failures that produced it rather than as an abstraction. SOLANGE computes real quantities over stand-in systems: a CAS(2e,2o)/STO-3G calculation on toluene where a phenylalanine sidechain belongs, on methanethiol where a cysteine does. Those are correct calculations of what they describe — two electrons of a model compound — and they are precisely what a pipeline demonstration requires. They become misconduct at the moment a report allows one to be read as a characterisation of a protein active site. The distance between we computed this and this is true of the target is the entire distance between a demonstration and a claim, and nothing about the arithmetic marks where that line falls.

Three observed failures fixed the principle's shape. A correlation energy computed over two electrons of a shared model compound was displayed beneath each target's classification; because those compounds are reused across genes, four unrelated targets showed an identical value and read as independent corroboration of one another — one calculation rendered four times, presenting as four measurements agreeing. A DMRG classification run over a six-electron fragment was allowed to reclassify a forty-four-electron target it had never characterised, silently promoting an infrastructure test into a scientific verdict. And a curated gene-level classification, reasoned for one specific structural mutant, was inherited by a splice variant of the same gene whose premise it does not share. In each case the computation was sound and the reporting was not.

The generalizable claim is therefore narrow and, the candidate argues, transferable: in any pipeline where cheap proxies stand in for expensive systems — which is to say most of computational science — scope cannot be left to the reader, because it will not survive contact

with a screenshot, an exported CSV, or a colleague reading a table six months later. It must be enforced structurally: grades computed from inputs rather than asserted by their producers, confidence bounded by the weakest input that determined a result, measurements authoritative only over the space they actually covered, and the qualification emitted from a single canonical definition into every surface a number can reach. The instantiation in SOLANGE is a set of mechanisms (evidence.py , the Gateway's provenance ceilings, the scope gates, a cross-source consistency checker that is run rather than trusted); the contribution is the vocabulary and the discipline, which are what would transfer to a platform sharing none of this code.

Applied to this dissertation's own results, DP5 is unambiguous. No target in this platform is presently validated to a standard supporting a chemical or biological finding, and every artifact the system emits says so in the payload rather than in a footnote. What has been independently established is narrower and is stated as such: the pipeline reproduces a reference it did not itself produce — N 2 CAS(6,6)/cc-pVDZ, recovered to 0.00 mHa classically and to 0.471 mHa (29% of chemical accuracy) through a twelve-qubit VQE on the Jordan–Wigner Hamiltonian. That validates the instrument for small-molecule bond breaking and transfers nothing to protein active sites: reproducing a known answer validates an instrument, not a discovery. Stating the limit alongside the result is not modesty about the work; under DP5 it is the work.

Positioning Against Precedent Systems — the White Space

The four properties that matter jointly are: cross-vendor/heterogeneous orchestration; signed, independently re-verifiable provenance (not merely a run log); the ability to cross a hard security boundary (interactive 2FA, no inbound) without weakening it; and an explicit regulatory-records mapping (21 CFR Part 11). No precedent system reviewed in §3.4 occupies the intersection of all four; that intersection is the defensible contribution.

System	Heterogeneous orchestration	Signed, re-verifiable provenance	Crosses hard security boundary (2FA)	Part-11 regulatory mapping
Airflow / Parsl / Covalent (workflow orchestrators)	✓	✗ (run logs, not sealed)	✗	✗
Qiskit Runtime / Orquestra (quantum orchestration)	▸ (single/allied vendor)	✗	✗	✗
W3C-PROV / scientific-provenance frameworks	n/a (schema, not orchestrator)	▸ (lineage; sealing optional)	✗	✗
Cleveland Clinic–IBM QCSC workflow (Mar 2026)	✓ (HPC ↔ QPU)	✗	✗	✗
SOLANGE (this work)	✓	✓ (LEON)	✓ (pull-based)	✓ (P1–P9)

✓ present · ▸ partial · ✗ absent. The claim is deliberately narrow: not "the first orchestrator," nor "the first provenance schema," but the first architecture to hold all four columns at once for oncology-mutation simulation under a regulatory mapping (§3.4, §06.ii).

Evaluation Map — Demonstrated vs. Pending

DSR treats demonstration as a legitimate but preliminary evaluation episode; a doctoral contribution requires evaluation beyond it. The table below is deliberately candid about which is which — the pending rows are the Phase 4 work, pre-registered on OSF before evaluation begins so acceptance thresholds cannot be retrofitted (§05).

Evaluation dimension	Status	Evidence / remaining gap
Feasibility	✓ demonstrated	End-to-end loop runs: NGS-driven queue → pull-agent → in-cluster execution → LEON notarization → display, on USC CARC Laguna.
Correctness / integrity	✓ demonstrated	Independent verification gates (cluster-side FCI vs CASSCF, Jordan-Wigner diagonalization, server-side seal + consistency re-check); no hard-coded results; monotonic energies across CAS sizes.
Comparison against baselines	► partial	Positioning argued qualitatively (table above); head-to-head evaluation against a push-based and an unsigned-provenance baseline is pending.
Utility with stakeholders	✗ pending	DSR Phase 4: structured evaluation with domain researchers and a regulatory-informed reviewer; OSF-preregistered acceptance criteria.
Generalization	► partial	Principles DP1–DP5 stated artifact-independently; a second-domain / second-workload instantiation to evidence transfer is designated future work.
Theoretical grounding	✓	Design Science Research paradigm (Hevner et al., 2004; Gregor & Hevner, 2013); provenance and zero-trust literatures.

The honest reading is that the artifact has crossed the feasibility and integrity thresholds convincingly, and that the remaining contribution work is evaluation, not construction. This is the expected posture at proposal stage under DSR: the design principles and their instantiation are established; their rigorous, stakeholder-facing evaluation is the dissertation's Phase 4 deliverable.

§ 07 Preliminary Work — A Deployed Prototype.

A fully functional prototype platform is deployed and accessible. All research artifacts are version-controlled on GitHub and directly accessible below.

7.1 • Deployed Artifacts & Research Outputs

- SOLANGE Platform - Live prototype + 3D simulation
- SOLANGE Platform Architecture
- SOLANGE Operational Stack

7.2 • Executed Classical HPC Runs — the Ladder's Classical Rung

Before the quantum rung, the classical rung of the escalation ladder (§06.ii) was exercised on real cluster hardware. Two exact active-space diagonalizations were dispatched from SOLANGE and executed on a Laguna (USC CARC) GPU node by the pull-agent — that is, they crossed the cluster's interactive-2FA boundary through the outbound-only agent channel exactly as DP2 specifies, and each returned a complete P1–P9 record sealed and re-verified by LEON. Both use the same TP53 C275F native model compound (methanethiol) as the anchor QPU runs of §7.4, so the two rungs are directly comparable.

Table 7.1 — Executed classical HPC runs (exact active-space diagonalization, USC CARC Laguna). Energies in Hartree. "Exact" means the ground state of the chosen active space is obtained by direct diagonalization — there is no variational approximation and hence no VQE-vs-exact gap to report; the computed value is the reference.

Date	Target	Active space	Qubits	Device	Wall-clock	Ground-state energy
26 Jul	TP53 C275F native	CAS(8e, 8o) / 6-31G	16	NVIDIA L40S · Laguna	100.5 s	−437.72232687
26 Jul	TP53 C275F native	CAS(10e, 10o) / 6-31G	20	NVIDIA L40S · Laguna	260.5 s	−437.75037505

Two properties of the ladder are visible directly in these two rows. First, the classical rung reaches widths the current quantum rung does not: 16 and 20 qubits here, against the fixed CAS(2e,2o) / 4-qubit scale of every executed QPU run (§7.4) — the classical rung is where larger active spaces are actually solved today. Second, the cost of that reach grows steeply: adding four qubits (16 → 20) lowered the ground-state energy by ≈ 28 mHa — more electron correlation captured by the larger active space — but at $2.6\times$ the wall-clock (100.5 s → 260.5 s). That super-linear growth, extrapolated past the ~ 18 -electron exact-classical wall (Table 8.1), is precisely what makes the eventual quantum rung necessary rather than merely available. As with §7.4, these are minimal proxy runs validating the pipeline and the ladder logic, not full-scale chemical-accuracy results for the 44-electron binding site (§08).

7.3 • Prototype Interface & Visualization

The prototype demonstrates the complete pipeline flow, grounded in the actual patient NGS report MI25-0349. Simulation results are scientifically informed illustrative values. The 3D visualization is a schematic representation of known TP53 secondary structure, loop topology, and zinc coordination — not crystallographic data. Both are sufficient to demonstrate the architecture's design logic and stakeholder communication capability.

7.4 • Executed Quantum-Hardware Runs — Pipeline Validation on Real QPUs

Six jobs have been executed on real IBM quantum hardware between 15 and 24 July 2026, on two processors — `ibm_marrakesh` (a Heron r2 device) and `ibm_kingston`. Each was sampled at 4,096 shots (one early run did not capture the shot count in its record), and each returned a complete P1–P9 provenance record populated from live device telemetry — calibration epoch (P3), per-gate and readout error,

and T_1/T_2 (P4) — sealed at source and re-verified by LEON on ingestion. These are deliberately scoped as pipeline-validation runs, not scientific results. Every run is a fixed-state expectation-value measurement — $\langle H \rangle$ evaluated on the Hartree-Fock reference state, not a ground-state search — at the minimal CAS(2e,2o) active space (4 qubits; 2 qubits for the H_2 control). What they establish is an information-systems result: that the full governed loop (mutation-aware Hamiltonian $\rightarrow$ real QPU execution $\rightarrow$ live-telemetry provenance capture $\rightarrow$ cryptographic seal $\rightarrow$ LEON notarization) runs end-to-end on real hardware across two devices and multiple targets, exactly as the architecture specifies.

Table 7.2 — Executed quantum-hardware runs (fixed-state $\langle H \rangle$, CAS(2e,2o)). Energies in Hartree. The reference is the exact $\langle H \rangle$ on the same Hartree-Fock state, computed classically; $|\Delta|$ is therefore the hardware deviation in reproducing a known expectation value — device noise, not an HF-vs-ground-state gap. The final row is the statevector control (see below).

Date	Target	Device	Qubits	Measured $\langle H \rangle$	Exact same-state ref.	$ \Delta $ (hardware noise)	Readout error
15 Jul	H_2 (control target)	ibm_marrakesh	2	-1.83694414	-1.83696586	2.2×10^{-5}	5.2×10^{-3}
15 Jul	TP53 C275F native	ibm_marrakesh	4	-432.90018127	-432.89297587	7.2×10^{-3}	3.2×10^{-3}
16 Jul	TP53 C275F mutant	ibm_kingston	4	-266.47261721	-266.47278261	1.7×10^{-4}	9.4×10^{-3}
19 Jul	TP53 C275F native	ibm_kingston	4	-432.89714609	-432.89297587	4.2×10^{-3}	5.9×10^{-3}
24 Jul	TP53 C275F mutant	ibm_kingston	4	-266.46986991	-266.47278261	2.9×10^{-3}	6.2×10^{-2}
24 Jul	KEAP1 LOF mutant	ibm_kingston	4	-432.89730045	-432.89297883	4.3×10^{-3}	5.3×10^{-2}
19 Jul	TP53 C275F native — statevector control	local simulator	4	-432.89297587	-432.89297587	0 (exact)	n/a

The control row carries the argument. Executed through the identical code path with the quantum device replaced by an exact statevector, the measurement reproduces the classical reference to the last recorded digit ($\Delta = 0$). The harness therefore contributes no error of its own, and the hardware deviations above — spanning 1.7×10^{-4} to 7.2×10^{-3} Ha across the two devices — are attributable to the QPUs (shot noise plus gate and readout error at this width), not to the pipeline. This is the intended Phase-3B-precursor evidence: the governance and orchestration layers hold when the backend is real quantum hardware returning genuinely probabilistic output, which is precisely the condition §06.iii's provenance schema and §06.ii's LEON notarization were designed for. The scientific targets — the ~24-electron local site and the ~44-electron / ~88-qubit full binding site of §1.2.1 — remain Phase 3B work (§08); nothing in this subsection is a chemical-accuracy claim. What is claimed, and what the records support, is that the pipeline executes and seals real quantum-hardware runs end-to-end.

§ 08 Limitations & Scope — What This Is, And Isn't.

Simulation Results

Binding energy values in the current prototype are scientifically informed estimates, not outputs of executed quantum simulations. Phase 3 proof-of-concept will produce real computed outputs on a simulator backend.

3D Visualization

Atomic coordinates are procedurally generated and schematic; they correctly represent structural topology but are not derived from PDB crystallographic data.

Hardware Scale & Classical Simulability

Full 44-electron pocket simulation on IBM Heron r3 is future work. The dissertation contribution is the architecture that enables this execution, not the execution itself. Full-scale quantum simulation of TP53 C275F (~44 electrons, ~88 qubits) on IBM Heron r3 is designated as a post-dissertation research phase, contingent on IBM Quantum Network institutional access. The architecture is hardware-agnostic at the workflow level and vendor-specific at the compilation and error-mitigation layer — IBM Heron r3 is one supported backend among several, with documented translation overhead at the lower layers.

Table 8.1 — State-vector memory-scaling rationale. The table below quantifies why classical simulation of quantum circuits becomes intractable as qubit count grows, and contextualises the Phase 3A scaling choices. A quantum circuit on n qubits requires a complex state-vector of length 2^n ; stored in double precision (complex128, 16 bytes per amplitude), the memory footprint is $2^n \times 16$ bytes. Every additional qubit doubles the requirement.

Qubits n	State-vector size	Memory (complex128)	Hardware regime / Phase 3 role
4	16 amplitudes	256 B	Laptop / cloud VM · Phase 3A minimal — live
30	$\sim 1 \times 10^9$	16 GB	Single GPU / workstation
40	$\sim 1.1 \times 10^{12}$	~ 17.6 TB	Multi-node HPC (NSF ACCESS / USC CARC) · Phase 3A intermediate — operational target ($\sim 20e/40q$)
45	$\sim 35 \times 10^{12}$	~ 563 TB	Leadership-class HPC (Frontier, Fugaku)
48	$\sim 281 \times 10^{12}$	~ 4.5 PB	Edge of global classical capacity · Phase 3A intermediate — scientific ceiling ($\sim 24e/48q$)
88	$\sim 3 \times 10^{26}$	$\sim 4.6 \times 10^9$ PB	Physically intractable on any classical machine · Phase 3B — requires IBM Heron r3 QPU ($\sim 44e$ full site)

Phase 3A intermediate targets $\sim 24e/48q$ as its scientific scale — the precise point where both exact classical diagonalisation (the determinant count turns combinatorial past $\sim 18e$) and classical state-vector simulation (~ 4.5 PB) reach their respective limits. For operational execution, the active space is reduced to $\sim 20e/40q$ (~ 17.6 TB) via an exact state-vector simulation of the quantum circuit (not a tensor-network approximation), which is within multi-node HPC reach on NSF ACCESS or USC CARC. This reduction does not compromise the architectural objective, which is operational stability rather than scientific validity: the pipeline is demonstrated at quantum-circuit scale two orders of magnitude beyond the minimal prototype, under full P1–P9 provenance governance. Phase 3B ($\sim 44e/88q$ on IBM Heron r3) exits the classically simulable regime entirely — the point at which the quantum algorithm is deliberately swapped from VQE to SqDRIFT, and the reason quantum hardware is required.

Production Deployment Model (Forward-Looking)

The proof-of-concept described in §07 is operated by a researcher who starts each pull-agent by hand in a terminal session — the correct workflow for a dissertation demonstration, but not the workflow a real deployment should require of its end user. The distinction that matters is who starts an agent and how, not what the end user experiences: even in the proof-of-concept, once an agent is running, the researcher interacts with SOLANGE exclusively through the browser. At production scale, that same separation is formalized — the

HPC/DMRG and QPU agents are installed once, by the customer's own IT or HPC administrator, as managed background services (a systemd unit or batch-scheduler-integrated job on the cluster; a systemd / launchd unit or container on a QPU-access host) that auto-start on boot, rather than a script a researcher re-runs by hand. The end user — a clinician or research scientist, not a systems engineer — never opens a terminal at any point.

This is not a departure from the pull-architecture defended in §06.ii (DP2): it is the same no-inbound-channel, no-externally-stored-credential guarantee, simply installed as a managed service instead of a manual script. Figure — below — makes the resulting trust boundary explicit: no arrow crosses from SOLANGE into the customer's HPC cluster or QPU-access host in either deployment mode; every result crosses that boundary only because an agent inside it chose, on its own authenticated session or service identity, to pull work and push a sealed result back out. Separating installation (IT, once) from use (clinician, always browser-only) is what makes the governance model of this dissertation credible as something that could ship into a regulated site, rather than only as a research demonstration. One half of this claim now has empirical grounding rather than remaining purely forward-looking: the classical rung's agent can run as an independent SLURM batch job (sbatch , requesting its own memory/time allocation) rather than only inside a researcher's interactive terminal session — exercised on USC CARC Laguna, surviving disconnection of the submitting terminal entirely, and self-limiting its own wall-clock budget to save a provisional result rather than losing a run outright if the SLURM allocation's time limit is reached first. This validates batch-scheduler integration specifically, one of the two installation mechanisms named above. What remains unvalidated, and is designated post-dissertation implementation work, is the fully unattended half of the claim: auto-start on boot with no sbatch / salloc invocation by any human at all, single sign-on integration, and disaster-recovery failover — none of which this dissertation's proof-of-concept exercises.

Intended Users — Present and Prospective

The prototype's interface language (§07) has historically framed its output for a clinician — "patient," "clinical dashboard." That framing states the platform's motivating problem honestly (§1.1) but does not describe who can act on the artifact's output today . Two DSR stakeholder classes do:

Computational-chemistry / drug-discovery researchers (academic or pharma R&D) are the platform's present users. The BQP-relevant classification framework (§06.i) answers a question this population asks directly — which non-druggable targets are worth a quantum run at all — and the governed pipeline (§06.ii) gives them a reproducible, re-verifiable result to build on. The interface now surfaces this classification as a headline count (Class A / Class B targets), not only as a table column, so the artifact's primary current utility is legible on first read.

Regulatory / compliance informatics functions within pharma are the platform's prospective users, and the reason its distinguishing contribution is a governance layer (§06.iii) rather than a faster energy calculation. A function evaluating whether quantum-derived molecular evidence could support a future IND submission needs the result to be auditable, reproducible, and cryptographically attestable — precisely what LEON and the P1–P9 record provide (§06.ii, §06.iv). The interface exposes this concretely as well: an exact, server-side count of LEON's append-only notarization events, not merely a description of the audit trail's existence. This user class is not empirically validated in this dissertation — quantum-derived evidence is not yet part of any regulatory submission pathway, and DSR Phase 4's expert-panel evaluation (§04, Layer 3) does not include a compliance-function participant. It is stated here as the artifact's intended future stakeholder, consistent with the forward-looking framing already applied to the Production Deployment Model above, not as a claim of present adoption.

The clinician remains the platform's long-run motivating persona (§1.1, §09) — the population whose treatment options this line of research ultimately aims to expand — but is explicitly not claimed as a present end user of the running artifact: the gap between a computed

binding energy and an actionable clinical recommendation is not bridged by this dissertation and is out of scope (§08, IND Regulatory Boundary).

Regulatory Scope

The compliance framework addresses data governance and audit trail requirements. It does not constitute an FDA submission or clinical trial regulation.

Generalizability

The platform is designed around NSCLC mutations but the orchestration architecture is mutation-agnostic and applicable to any BQP-relevant molecular simulation target.

Single-Modality Anchor

The proof-of-concept execution anchors on superconducting qubits (§1.4). This decision is the best-supported choice at the time of submission, but it narrows the empirical claims of the dissertation. Multi-modality validation — re-running the pipeline against trapped-ion or neutral-atom backends to demonstrate behavioral equivalence — is deferred to future work. The architecture is described abstractly enough that re-anchoring on a different modality is a defined exercise should the empirical landscape shift before defense.

Nvidia Substrate (Forward-Compatible, Not Empirically Exercised)

The post-April-2026 hybrid substrate — Nvidia CUDA-Q , NVQLink , and Ising — is acknowledged in §1.4.1 and treated as commodity infrastructure in the three-layer architecture (§06.ii). Phase 3 of this dissertation does not execute on that substrate. IBM is absent from Nvidia's announced NVQLink partner roster (Atom Computing, IonQ, IQM, Q-CTRL, Academia Sinica, EeroQ), and the SqDRIFT empirical anchor ($C_{13}Cl_2$, IBM Research 2026; arXiv:2603.08696) runs on IBM Heron via Qiskit Runtime with deterministic decoding (pyMatching). The contribution is therefore structural , not empirical: (a) the orchestration architecture is designed to interoperate with the Nvidia substrate (§06.ii property 3, bifurcation-aware), but Phase 3 exercises only the IBM-native side; (b) the ninth element P9 of the Part-11 provenance record (§06.iii) is forward-compatible — it specifies what records a regulated workflow will require when ML-decoded QEC enters production, but Phase 3's P9 records reduce to the minimal deterministic form (algorithm name + software version). Empirical validation of bifurcation-aware orchestration across both stacks, and of ML-decoder provenance on a live Ising-class decoder, are designated post-dissertation research phases.

IND Regulatory Boundary

The compliance framework (§06.iii) addresses the information-systems layer of FDA 21 CFR Part 11 — what records must be retained, in what form, against what validation criteria. It does not address the broader regulatory-science questions: the criteria under which quantum-simulated binding affinities could be admitted as evidence in an Investigational New Drug (IND) submission, clinical-trial protocol design implications of using quantum-derived candidates, or FDA's separate guidances on Model-Informed Drug Development. These topics fall outside the dissertation's scope and are reserved for future regulatory-science work.

§ 09 Significance — Why This Matters.

Of the 1.53 million patients who die annually from NSCLC, a significant population carries non-druggable mutations for whom no computational drug discovery pathway currently exists. Quantum hardware has reached sufficient capacity. The algorithms exist. The clinical data exists. What is missing is the information systems infrastructure that connects them — governed, reproducible, and accessible to the cross-disciplinary stakeholders who need to act on the results. The platform is designed around four NSCLC anchor mutations but the architecture is extensible to over one million patients annually across fifteen pipeline-eligible non-druggable mutations — the largest unaddressed population in computational oncology within the quantum feasibility window.

This dissertation builds that infrastructure. Not as a physics paper claiming a new binding energy, but as an IT systems contribution claiming a new architecture. The distinction is precise and important:

Translational Pipeline — Where the Architecture Sits in the Larger Pathway

This subsection situates the IT contribution inside the wider drug-development arc; the downstream stages (in vitro through Phase III) are not claimed as dissertation contributions. The architecture's role is bounded to the In Silico stage — providing governed, reproducible, audit-traceable computational outputs that downstream stages can cite as a primary record. The clinical and regulatory stages that follow are shown to make the IT contribution's significance concrete, not to extend the dissertation's scope into pharmacology or clinical trials.

This Dissertation

Quantum simulation of C275F local (~24e) and full (~44e) binding sites using the QC·AI·HPC orchestration platform. Output: ranked drug candidate molecules with quantum-accurate binding energies, governed by FDA 21 CFR Part 11-compliant audit trails.

Biochemical Validation

p53 binding assays, transcriptional activation assays, and cell-free DNA binding experiments on shortlisted candidates. Confirms that computationally predicted binding translates to measurable biochemical activity.

Animal Model Studies

Efficacy and preliminary safety profiling in murine lung cancer models harboring the C275F mutation. Provides pharmacokinetic and toxicity data required before human studies.

Human Safety & Dosing

First-in-human studies establishing safety profile, maximum tolerated dose, and pharmacokinetics in cancer patients.

Efficacy Signals

Targeted enrollment of TP53 C275F-positive NSCLC patients. Primary endpoint: objective response rate. First clinical evidence that quantum-designed candidates produce measurable anti-tumor activity.

Regulatory Approval

Randomized controlled trials against standard of care. Successful completion enables FDA submission and clinical deployment — completing the translational arc from quantum simulation to approved therapy for a currently non-druggable mutation.

This framing positions the quantum simulation architecture not as an isolated computational exercise, but as the critical first step in a clinically meaningful and scientifically grounded drug discovery pipeline targeting currently non-druggable TP53 mutations in lung cancer.

Future Research Directions

The molecular-level simulations enabled by this dissertation's SqDRIFT+VQE pipeline are positioned to inform a second clinical question that lies beyond the dissertation's scope: NGS-informed radiotherapy decision support. The non-druggable mutations targeted here — TP53, KEAP1, CDKN2A, STK11 — are the same mutations most frequently implicated in radioresistance literature, and the quantum-derived mechanistic features they expose (DNA damage repair dynamics, oxidative-stress response via the NRF2 pathway disrupted by KEAP1 mutations, apoptotic signaling) are precisely the features current radiogenomic models lack. The signal is measurable in contemporary clinical cohorts: in a stage III NSCLC cohort treated with definitive chemoradiation and consolidative durvalumab, tumor mutational burden was the covariate significantly associated with 24-month local-regional failure and progression-free survival [64] — the class of clinical endpoint the radiogenomic layer would predict. Bridging quantum-simulated mutation effects into radiotherapy planning — extending prior efforts such as the Genomic-Adjusted Radiation Dose (GARD) framework (Scott et al., 2017) — is a defensible post-doctoral research program that this dissertation enables but does not undertake. The dissertation builds the molecular-simulation infrastructure; the radiotherapy decision-support layer, with its independent regulatory pathway and prospective-validation requirements, is reserved for future work.

From Governed Pipeline to Chemist-Operable Instrument


The three claims of this dissertation (§06) are about the orchestration architecture — that a fermionic-simulation pipeline can be routed, notarized, and access-controlled across organizational and hardware boundaries while remaining auditable under an FDA Part 11 reinterpretation. Every claim is demonstrated end to end in the deployed prototype (§07) and is independent of any specific chemical result the pipeline happens to compute. Making that same prototype operable by an industrial computational chemist against a real binding site — rather than against the CAS(2e,2o)/STO-3G proxy compounds used throughout Phase 3A to exercise the architecture (§07, §08) — is a separate body of work, and this dissertation is explicit about which side of that line each remaining gap falls on.

Two gaps were engineering-completable and within the Information Systems contribution's natural extension, and both have since been closed in the deployed prototype. The first was accepting a real protein structure (PDB or AlphaFold model) as input and carving an actual quantum-mechanical cluster around the mutated residue, in place of the small stand-in molecules (e.g., methanethiol, toluene, formamide — 24 to 50 total electrons) previously mapped one-per-gene. A target now resolves to its deposited structure, the structure is retrieved from the RCSB PDB, and a cluster is carved around the mutated residue with valence capping, residue-level charge assignment, and a detector for missing hydrogens; a 195-atom cluster around TP53 Cys275, carved from 2OCJ, has been built and carried through a converged mean-field solution on the Laguna cluster. The second was making the orchestration layer's active-space parameters (Computational Decision Gateway, §06.i; dispatch controls, §07) size-aware per target rather than fixed. The dispatch surface now rebuilds its active-space selector against the selected target's own active-electron count, bounds the offered sizes at the exact-diagonalisation orbital wall, and escalates the solver automatically beyond that wall, so a researcher is no longer offered a control that cannot succeed for the compound behind it.

A third gap was not an engineering gap at all: the DMRG classical-classifier stage (Rung 3 of the orchestration ladder, §07) performed its active-orbital selection via conventional CASSCF, whose full-configuration-interaction step is combinatorial in the number of active orbitals and becomes impractical past roughly sixteen — the same ceiling that bounds the exact-classical and VQE rungs (§06.i, §07). The deployed prototype therefore did not realise DMRG's own methodological premise, that a large active space can be handled at polynomial rather than exponential cost. That gap has since been closed, and it was the one item of the three anticipated to require computational-chemistry method development rather than software architecture. The orbital optimisation is now performed by the DMRG sweep itself: a

tensor-network solver replaces the full-configuration-interaction solver inside the self-consistent-field loop, so the reachable active space is no longer bounded by the determinant count. The implementation is validated rather than asserted — its one- and two-particle reduced density matrices are compared element by element against exact full configuration interaction on a small random Hamiltonian before every optimisation, and every solver call independently reconstructs its own energy from those matrices and checks it against the sweep — and it has been exercised at CAS(20,20), forty qubits, past the wall that bounds the exact-classical and VQE rungs. The boundary drawn in §08 is unchanged by this: what the implementation establishes is that the governed pipeline carries a larger active space under the same provenance and notarisation, not any chemical claim about the compound it carried. The architecture remains the artifact; the chemistry remains outside the contribution. What remains open is accordingly not the method but the measurement. The Gateway's evidence-based scope gate (§06.i — a measurement counts as evidence only if it covers at least 90% of a target's active site) has not yet been satisfied for any of this dissertation's real oncology targets: the memory ceiling that blocked it has since been lifted, and removing it revealed that it was not the binding constraint. The ~44-electron figure states what a full treatment of the binding pocket requires; it does not specify which orbitals those are, and no such specification exists — the active-space selections the pipeline can actually execute (sulfur 3p, or zinc 3d with sulfur 3p) span a few orbitals rather than forty-four. The run is therefore not waiting on compute. It is waiting on a chemically motivated definition of the pocket's active space, which is the chemist's input under the boundary drawn in §08 and not something the orchestration layer can supply. Until such a run completes, the Gateway will correctly and consistently decline to certify any of these targets as quantum-justified on measured grounds, which is the system behaving exactly as designed rather than a defect to route around, and C275F's Class B designation stands as recorded in §06.i — resting on the state of the published evidence rather than on a measurement of its own site

§ 10 Proposed Timeline — The Road Ahead.



[Figure — rendered from source]

Phase	Activity	Status	Target
1	Foundation and literature review	✓ Complete	—
2	Platform prototype — deployed	✓ Complete	—
3A	Architecture validation on PennyLane simulator — minimal: 4-qubit CAS(2e,2o) ✓ live; intermediate: local site scientific target ~24e/48q (classical-chemistry ceiling + state-vector memory limit ~4.5 PB); operational run ~20e/40q (~17 TB, within HPC reach) ► pending	✓ Partial	May–Aug 2026
3B	Architecture scaling run — full site (~44e) on IBM Heron r3 (contingent on IBM Quantum Network access)	Contingent	Sep–Dec 2026
4	Triangulated DSR evaluation — Layers 1–4 across sub-phases 4A–4F (technical, comparative, expert panel n=4–6, case-based walkthrough)	Pending	Jan–Apr 2027
5	Dissertation writing (chapters 1–8)	Pending	May–Sep 2027
6	Pre-defense committee revisions	Pending	Oct 2027
7	Defense	Pending	Nov–Dec 2027
—	IBM Heron r3 full-scale runs (post-defense)	Future work	2028+

Timeline Rationale & Contingencies

The timeline allocates ~20 months from proposal acceptance to defense — a deliberate revision from earlier drafts that compressed the schedule into 9 months. The realistic distribution recognizes that (a) Phase 3A is a four-month engineering effort, not a two-week sprint; (b) Phase 3B is contingent on IBM Quantum Network institutional access, which can take 3–6 months to secure; (c) recruiting 4–6 cross-disciplinary domain experts and running structured DSR evaluation cannot be compressed below one quarter without losing methodological rigor; and (d) writing a defensible IT/DSR dissertation while balancing professional commitments requires a sustained five-month drafting window plus a one-month revision cycle.

Phase 3B is explicitly designated as contingent. The dissertation's central contribution — the orchestration architecture and the FDA Part 11 reinterpretation — is independent of full-pocket execution on physical quantum hardware. If IBM Quantum Network access is not granted within the Sep–Dec 2026 window, Phase 3B is documented as future work in the limitations chapter, and the defense proceeds on

the strength of Phase 3A (architecture validation) and Phase 4 (expert panel evaluation). This contingency is the practical expression of the "architecture is the contribution; simulation is the application" framing.

§ 11 Selected References — The Shoulders.

- 01 Abrams, D. S., & Lloyd, S. (1999). Quantum algorithm providing exponential speed increase for finding eigenvalues and eigenvectors. *Physical Review Letters* , 83(24), 5162–5165.
- 02 Aspuru-Guzik, A., Dutoi, A. D., Love, P. J., & Head-Gordon, M. (2005). Simulated quantum computation of molecular energies. *Science* , 309(5741), 1704–1707.
- 03 Bauer, B. et al. (2020). Quantum algorithms for quantum chemistry and quantum materials science. *Chemical Reviews* , 120(22), 12685–12717.
- 04 Bernstein, E., & Vazirani, U. (1997). Quantum complexity theory. *SIAM Journal on Computing* , 26(5), 1411–1473.
- 05 Cao, Y. et al. (2019). Quantum chemistry in the age of quantum computing. *Chemical Reviews* , 119(19), 10856–10915.
- 06 Carleo, G., & Troyer, M. (2017). Solving the quantum many-body problem with artificial neural networks. *Science* , 355(6325), 602–606.
- 07 FDA (2003). 21 CFR Part 11: Electronic records; electronic signatures — guidance for industry. U.S. Food and Drug Administration.
- 08 Hevner, A. R., March, S. T., Park, J., & Ram, S. (2004). Design science in information systems research. *MIS Quarterly* , 28(1), 75–105.
- 09 Joerger, A. C., & Fersht, A. R. (2016). The p53 pathway: Origins, inactivation in cancer, and emerging therapeutic approaches. *Annual Review of Biochemistry* , 85, 375–404.
- 10 McClean, J. R. et al. (2016). The theory of variational hybrid quantum-classical algorithms. *New Journal of Physics* , 18(2), 023023.
- 11 Peffers, K., Tuunanen, T., Rothenberger, M. A., & Chatterjee, S. (2007). A design science research methodology for information systems research. *Journal of Management Information Systems* , 24(3), 45–77.
- 12 Peruzzo, A. et al. (2014). A variational eigenvalue solver on a photonic chip. *Nature Communications* , 5, 4213.
- 13 Santagati, R. et al. (2024). Drug design on quantum computers. *Nature Physics* , 20, 549–557.
- 14 Sarkar, I. N. (2010). Biomedical informatics and translational medicine. *Journal of Translational Medicine* , 8, 22.
- 15 Scott, J. G., Berglund, A., Schell, M. J., Mihaylov, I., Fulp, W. J., Yue, B., Welsh, E., Caudell, J. J., Ahmed, K., Strom, T. S., Mellon, E., Venkat, P., Johnstone, P., Foekens, J., Lee, J., Moros, E., Dalton, W. S., Eschrich, S. A., McLeod, H., Harrison, L. B., & Torres-Roca, J. F. (2017). A genome-based model for adjusting radiotherapy dose (GARD): A retrospective, cohort-based study. *The Lancet Oncology* , 18(2), 202–211.
- 16 Venable, J., Pries-Heje, J., & Baskerville, R. (2016). FEDS: A framework for evaluation in design science research. *European Journal of Information Systems* , 25(1), 77–89.
- 17 Yin, R. K. (2014). Case study research: Design and methods (5th ed.). Sage Publications.

Genomic Databases & Epidemiological Sources

- 18 Tate, J. G. et al. (2019). COSMIC: The Catalogue of Somatic Mutations in Cancer. *Nucleic Acids Research* , 47(D1), D941–D947. cancer.sanger.ac.uk/cosmic
- 19 Bouaoun, L. et al. (2016). TP53 Variations in Human Cancers: New Lessons from the IARC TP53 Database and Genomics Data. *Human Mutation* , 37(9), 865–876. tp53.isb-cgc.org
- 20 The Cancer Genome Atlas Research Network (2014). Comprehensive molecular profiling of lung adenocarcinoma. *Nature* , 511, 543–550. Data portal: portal.gdc.cancer.gov
- 21 Bray, F. et al. (2024). Global cancer statistics 2022: GLOBOCAN estimates of incidence and mortality worldwide for 36 cancers in 185 countries. *CA: A Cancer Journal for Clinicians* , 74(3), 229–263.
- 22 Joerger, A. C., & Fersht, A. R. (2007). Structural biology of the tumor suppressor p53 and cancer-associated mutants. *Advances in Cancer Research* , 97, 1–23. (Cited in Tables 1–2 for TP53 structural mutant classification: R175H, R282W.)

Industry & Technical Sources

- 23 Drury, K. (2026, April 26). Nvidia just announced a major quantum computing development . *The Motley Fool*. [fool.com](https://www.fool.com)
- 24 FDA, Health Canada, & MHRA (2021). Good Machine Learning Practice for medical device development: Guiding principles . U.S. Food and Drug Administration.
- 25 IBM (2026, March 5). IBM and university researchers create a never-before-seen molecule and prove its exotic nature with quantum computing . *IBM Newsroom*. newsroom.ibm.com
- 26 NVIDIA (2025, October 28). NVIDIA introduces NVQLink — connecting quantum and GPU computing for 17 quantum builders and nine scientific labs . *NVIDIA Newsroom*. nvidianews.nvidia.com
- 27 NVIDIA (2026, April 14). NVIDIA launches Ising, the world's first open AI models to accelerate the path to useful quantum computers . *NVIDIA Newsroom*. nvidianews.nvidia.com

- 28 Quantinuum (n.d.). Quantinuum–NVIDIA partnership: ADAPT-GQE applied to imipramine . Quantinuum corporate blog. quantinuum.com
- 29 The Quantum Insider (2025, June 9). IonQ speeds quantum-accelerated drug-development application in partnership with AstraZeneca, AWS, and Nvidia . thequantuminsider.com
- 30 Swayne, M. (2026, March 26). Quantinuum researchers demonstrate quantum computations with dozens of protected logical qubits . The Quantum Insider. thequantuminsider.com
- 31 Quantum Computing Report (2026, March 26). Cleveland Clinic and IBM implement quantum workflow for protein simulation . Quantum Computing Report. quantumcomputingreport.com

Theoretical Foundations — Classical → Quantum Continuum & Enterprise Architecture

- 32 Alexeev, Y. et al. (2024). Quantum-centric supercomputing for materials science: A perspective on challenges and future directions. *Future Generation Computer Systems* , 160, 666–710. doi:10.1016/j.future.2024.04.060
- 33 Campbell, E. (2019). Random compiler for fast Hamiltonian simulation. *Physical Review Letters* , 123(7), 070503. doi:10.1103/PhysRevLett.123.070503
- 34 Feynman, R. P. (1982). Simulating physics with computers. *International Journal of Theoretical Physics* , 21(6/7), 467–488. doi:10.1007/BF02650179
- 35 Fowler, A. G., Mariantoni, M., Martinis, J. M., & Cleland, A. N. (2012). Surface codes: Towards practical large-scale quantum computation. *Physical Review A* , 86(3), 032324. doi:10.1103/PhysRevA.86.032324
- 36 Helgaker, T., Jørgensen, P., & Olsen, J. (2000). *Molecular Electronic-Structure Theory* . John Wiley & Sons. doi:10.1002/9781119019572
- 37 Kandala, A. et al. (2017). Hardware-efficient variational quantum eigensolver for small molecules and quantum magnets. *Nature* , 549(7671), 242–246. doi:10.1038/nature23879
- 38 Kiss, O., Grossi, M., & Roggero, A. (2023). Importance sampling for stochastic quantum simulations. *Quantum* , 7, 977. doi:10.22331/q-2023-04-13-977
- 39 Lankhorst, M. (2017). *Enterprise Architecture at Work: Modelling, Communication and Analysis* (4th ed.). Springer. doi:10.1007/978-3-662-53933-0
- 40 Levine, B. G., Durden, A. S., Esch, M. P., Liang, F., & Shu, Y. (2021). CAS without SCF—Why to use CASCI and where to get the orbitals. *The Journal of Chemical Physics* , 154(9), 090902. doi:10.1063/5.0042147
- 41 Lloyd, S. (1996). Universal quantum simulators. *Science* , 273(5278), 1073–1078. doi:10.1126/science.273.5278.1073
- 42 McArdle, S., Endo, S., Aspuru-Guzik, A., Benjamin, S. C., & Yuan, X. (2020). Quantum computational chemistry. *Reviews of Modern Physics* , 92(1), 015003. doi:10.1103/RevModPhys.92.015003
- 43 The Open Group (2022). *ArchiMate 3.2 Specification* (Document No. C226). The Open Group.
- 44 The Open Group (2022). *The TOGAF Standard, 10th Edition* (Architecture Development Method). The Open Group. ISBN 978-94-018-0859-0.
- 45 Pierantoni, G., Kiss, T., Bolotov, A. et al. (2023). Toward a reference architecture based science gateway framework with embedded e-learning support. *Concurrency and Computation: Practice and Experience* , 35(18), e6872. doi:10.1002/cpe.6872
- 46 Preskill, J. (2018). Quantum computing in the NISQ era and beyond. *Quantum* , 2, 79. doi:10.22331/q-2018-08-06-79
- 47 Reiher, M., Wiebe, N., Svore, K. M., Wecker, D., & Troyer, M. (2017). Elucidating reaction mechanisms on quantum computers. *Proceedings of the National Academy of Sciences* , 114(29), 7555–7560. doi:10.1073/pnas.1619152114
- 48 Sandve, G. K., Nekrutenko, A., Taylor, J., & Hovig, E. (2013). Ten simple rules for reproducible computational research. *PLOS Computational Biology* , 9(10), e1003285. doi:10.1371/journal.pcbi.1003285
- 49 Stodden, V. et al. (2016). Enhancing reproducibility for computational methods. *Science* , 354(6317), 1240–1241. doi:10.1126/science.aah6168
- 50 Tilly, J. et al. (2022). The variational quantum eigensolver: A review of methods and best practices. *Physics Reports* , 986, 1–128. doi:10.1016/j.physrep.2022.08.003
- 51 Wilkinson, M. D. et al. (2016). The FAIR Guiding Principles for scientific data management and stewardship. *Scientific Data* , 3, 160018. doi:10.1038/sdata.2016.18
- 52 W3C (2013). *PROV-O: The PROV Ontology* (W3C Recommendation, 30 April 2013). World Wide Web Consortium. w3.org/TR/prov-o
- 53 Zachman, J. A. (1987). A framework for information systems architecture. *IBM Systems Journal* , 26(3), 276–292. doi:10.1147/sj.263.0276
- 54 da Silva, M. P., et al. (2024). Demonstration of logical qubits and repeated error correction with better-than-physical error rates. *Nature* , 638, 920–926. doi:10.1038/s41586-024-08005-8 — Microsoft & Quantinuum: 800× logical error reduction, 12 logical qubits, 14,000 circuits without error on Quantinuum H2 trapped-ion processor.
- 55 Ransford, A., et al. (Quantinuum, in collaboration with Sandia National Laboratories) (2026). Helios: a 98-qubit trapped-ion quantum processor with high-fidelity all-to-all connectivity. *Nature* . doi:10.1038/s41586-026-10676-4 — 98 physical qubits (137 Ba +),

- 99.92% two-qubit gate fidelity, random-circuit-sampling quantum advantage; ~55 ms per depth-1 layer (throughput data point for §1.4.2).
- 56 Infleqtion, Inc. (2026, February 17). Infleqtion and Churchill Capital Corp X complete business combination [Press release]. Investor Relations. ir.infleqtion.com — First neutral-atom quantum company to list publicly (NYSE: INFQ; ~\$551M gross proceeds via business combination with Churchill Capital Corp X); systems deployed with the U.S. Department of War, NASA, DARPA, and the U.S. Department of Energy, and the U.K. National Quantum Computing Centre; NVIDIA joint logical-qubit materials-science demonstration (Anderson Impurity Model); roadmap target of 100+ logical qubits by 2028 on the Sqaile platform.
 - 57 Lee, S., Lee, J., Zhai, H., Tong, Y., Dalzell, A. M., Kumar, A., Helms, P., Gray, J., Cui, Z.-H., Liu, W., Kastoryano, M., Babbush, R., Preskill, J., Reichman, D. R., Campbell, E. T., Valeev, E. F., Lin, L., & Chan, G. K.-L. (2023). Evaluating the evidence for exponential quantum advantage in ground-state quantum chemistry. *Nature Communications* , 14(1), 1952. doi:10.1038/s41467-023-37587-6 — Cited in §1.2 as the counter-argument the BQP-relevance framing is explicitly delimited against: evidence for exponential quantum advantage across chemical space has not been established in the general case.
 - 58 IBM Quantum. (2026). Processor types . IBM Quantum Documentation. quantum.cloud.ibm.com/docs/en/guides/processor-types — Heron r3: 156-qubit heavy-hex QPU; used to substantiate the r3 fidelity/coherence characterisation in §1.4.
 - 59 IBM Quantum. (2026, January 5). Nighthawk and the latest Heron are now available . IBM Quantum product updates. quantum.cloud.ibm.com/announcements/en/product-updates/2026-01-05-nighthawk — Release of Heron r3 (ibm_pittsburgh), ~350 μ s T2 and improved gate/readout fidelity; source for the r3-vs-r2 fidelity argument in §1.4.
 - 60 IBM (2026, July 23). IBM to acquire HRL Laboratories to power the future of quantum . IBM Newsroom. newsroom.ibm.com — IBM acquires HRL Laboratories (formerly owned by Boeing and General Motors), adding silicon-spin qubit expertise and quantum sensing capabilities alongside IBM's existing superconducting qubit roadmap; transaction expected to close Q3 2026. Establishes IBM's formal two-track qubit strategy: superconducting (Heron r3 → Starling → Blue Jay) + silicon-spin (HRL/Anderson foundry, fault-tolerant regime ~2030+).
 - 61 Nellis, S. (2026, July 23). IBM buys HRL Laboratories to shift to two-track quantum computing strategy . Reuters. reuters.com — Reports IBM acquisition of HRL Laboratories; confirms two-track strategy (superconducting + silicon-spin); notes Boeing and GM retain quantum application partnerships post-acquisition; references the Anderson quantum wafer foundry established by IBM in 2026 with U.S. Department of Commerce support.
 - 62 Ross, J. S., Thummalapalli, R., Febres-Aldana, C. A., et al. (2026). Clinical significance of MTAP deletions and their overlap with concurrent oncogenic driver alterations including EGFR in NSCLC. *Journal of Thoracic Oncology* , 21(3), 103520. doi.org/10.1016/j.jtho.2025.11.010 — Across 4,926 NSCLC tumours profiled by MSK-IMPACT, 10% carried an MTAP deletion, which co-occurs with CDKN2A deletion; among stage-IV MTAP-deleted cases, 83% also carried an oncogenic driver (48% EGFR). Establishes the contemporary evidence base for CDKN2A-locus loss as a clinically consequential, indirectly addressable (synthetic-lethal, via MTA-cooperative PRMT5 inhibition) target rather than a directly druggable one — the pattern this dissertation's classification framework formalises. (PMID 41260457.)
 - 63 Volta, F., La Monica, S., Leonetti, A., et al. (2023). Intrinsic resistance to osimertinib in EGFR-mutated NSCLC cell lines induced by alteration in cell-cycle regulators. *Targeted Oncology* , 18(6), 953–964. doi.org/10.1007/s11523-023-01005-0 — Reduced CDKN2A (p16) with elevated Cyclin D1 confers intrinsic osimertinib resistance via the CDK4/6–Rb axis in EGFR-mutant NSCLC cell lines. Empirical support for the CDK4/6-interface rationale recorded for CDKN2A in Table 1.
 - 64 Lebow, E. S., Shepherd, A., Eichholz, J. E., et al. (2023). Analysis of tumor mutational burden, progression-free survival, and local-regional control in patients with locally advanced non-small cell lung cancer treated with chemoradiation and durvalumab. *JAMA Network Open* , 6(1), e2249591. doi.org/10.1001/jamanetworkopen.2022.49591 — In 81 patients, tumor mutational burden was the covariate significantly associated with 24-month local-regional failure and progression-free survival on multivariable analysis; KEAP1/NFE2L2-altered tumours showed higher local-regional failure rates but this was not the significant multivariate predictor. Contemporary evidence that mutation-burden signal in the mutation classes targeted here carries radiotherapy-response information, supporting the radiogenomic extension identified in §09 as future work.

Appendix A Complete Glossary — Ninety-Five Entries.

A Translational QC·AI·HPC Orchestration Platform for Non-druggable NSCLC Tumor-Suppressor Mutations

This glossary consolidates every acronym and technical term that appears across the dissertation proposal, conference poster, deployed prototype (solange-platform.bio), and supporting coursework artifacts. Each entry pairs the standard definition with the term's specific role in this dissertation, so a reader unfamiliar with quantum computing, computational chemistry, FDA regulation, or oncology can follow the work without external reference. Terms are listed alphabetically.

Total entries: 95

Acronym	Full Name	Definition & Context in This Dissertation
21 CFR	Title 21, Code of Federal Regulations	The U.S. federal regulations governing food and drugs, administered by the FDA. Two parts are relevant: Part 11 (electronic records - in dissertation scope) and Part 312 (IND submission - explicitly out of scope).
3D CNN	Three-Dimensional Convolutional Neural Network	A CNN extended to three-dimensional input grids. Architecture of the Nvidia Ising QEC decoder, applied to syndrome lattices in real-time error correction. Layer 2 commodity; provenance recorded in P9.
ACCESS	Advanced Cyberinfrastructure Coordination Ecosystem: Services & Support (NSF)	The U.S. National Science Foundation's national programme allocating time on academic supercomputers. Named alongside USC CARC as the classical HPC substrate on which the Phase 3A state-vector run (~20e/40q, ~17.6 TB) is within multi-node reach (§08).
ALK	Anaplastic Lymphoma Kinase	An oncogene whose rearrangements drive ~5% of NSCLC cases. ALK has FDA-approved targeted therapies (crizotinib, alectinib) and therefore represents the druggable counterpart to the tumor-suppressor mutations addressed in this dissertation. Cited in the Problem Statement to delineate the clinical gap.
AlphaFold	DeepMind protein-structure prediction system	Source of a predicted structure where no experimental one exists. Supplies the STK11 F354L local site (UniProt Q15831), and is named alongside experimental PDB entries as an admissible structure input for carving a quantum-mechanical cluster around a mutated residue (§09).
API	Application Programming Interface	A defined contract through which software components communicate. Referenced in problem statement P1 (no standardized handoff between classical and quantum compute) - the absence of a stable API contract between HPC pre-screening and quantum circuit builders is one of the gaps the orchestration platform addresses.
ATP	Adenosine Triphosphate	The principal energy-carrying molecule in cells; binds to kinase active sites. Relevant to STK11 - one of the four target mutations - whose kinase domain coordinates ATP and contributes to the active-space electron count requiring quantum-grade simulation.
AVAS	Atomic Valence Active Space	An automated procedure for choosing which orbitals enter a correlated calculation, by projecting the mean-field orbitals onto selected atomic valence functions. Used on the DMRG rung to derive the active space around the mutated residue rather than fixing it by hand.
B3LYP	Becke 3-parameter Lee-Yang-Parr (DFT functional)	A widely used hybrid density functional in classical computational chemistry. The poster's results section reports B3LYP at -8.4 kcal/mol with 2.8 kcal/mol error - outside the chemical accuracy threshold of ≤2 kcal/mol - establishing why classical DFT is insufficient for C275F at the active-space boundary.
BAA	Business Associate Agreement	A HIPAA-mandated contract between a covered entity (hospital, clinic) and any external party that handles Protected Health Information on its behalf. Out of scope for this dissertation because the platform uses fully de-identified single-case reference data (MI25-0349) but would become relevant in any future multi-patient extension.
bond dimension	M — the rank retained in a matrix-product-state representation	The single parameter controlling DMRG's accuracy and cost. It is bounded above by e S with the entanglement entropy S across the most strongly correlated bipartition of the active space, while cost scales as M ³ . The bound assumes a flat entanglement spectrum; measured on an N ₂ stretch series the growth exponent is 0.31, not 1. This is why entanglement, not electron count, decides whether a target is classically tractable — the physical basis of the §06.i classification.
BPR	Business Process Reengineering	A management methodology for analyzing and redesigning organizational workflows. Used in the IST 697 coursework deliverable that mapped the QC+AI+HPC drug discovery platform across five business processes (P1–P5), identifying the Molecular Simulation Pipeline as the main process where the BQP-class computational bottleneck concentrates.
BQP	Bounded-error Quantum Polynomial time	The complexity class of decision problems efficiently solvable on a quantum computer with bounded error probability. The dissertation's first original contribution (§06.i) is a BQP-relevance classification framework arguing that the fermionic Hamiltonian eigenvalue problem at C275F's 44-electron active-space boundary admits a known BQP-class quantum algorithm while best-known classical methods fail at chemical accuracy. This is an operational BQP-relevance claim — conditioned on the current classical state of the art and on sufficient initial-state overlap — not a formal classical-quantum separation, which remains open (ground-state finding is QMA-hard in general; see §1.2).
C ₁₃ Cl ₂	13-Carbon, 2-Chlorine half-Möbius molecule	The synthetic molecule used by IBM in its 2026 demonstration (IBM Research; arXiv:2603.08696) to run SqDRIFT at 32 electrons / 100 qubits - the empirical anchor of the dissertation's Layer 1 scientific claim. The transition from C ₁₃ Cl ₂ (synthetic, vacuum, cryogenic) to a clinically relevant metalloprotein mutation is explicitly framed as Layer 1 of the dissertation's novelty argument.
C275F	TP53 p.Cys275Phe (cysteine-to-phenylalanine substitution at position 275)	A non-druggable TP53 missense mutation in which Cys275 is replaced by Phe. The mutation introduces a rigid aromatic π-system into the hydrophobic β-sandwich core of the p53 DNA-binding domain, disrupting DNA-contact geometry without creating an exploitable drug-binding cavity. C275F is the dissertation's primary demonstration case - sourced from patient NGS report MI25-0349 - and the focal point of the Phase 3A and 3B simulations. C275F is classified as a structural mutant: the Cys → Phe substitution destabilizes the hydrophobic core of the DNA-binding domain rather than directly abolishing a DNA-phosphate contact. No crystallographic structure of C275F exists in PDB, consistent with the conformational heterogeneity expected from a structural mutation.
CARC	USC Center for Advanced Research Computing	The University of Southern California's research-computing centre and operator of the Laguna cluster, on which this dissertation's classical HPC and DMRG rungs execute. Its mandatory interactive multi-factor authentication is the concrete security posture that motivates DP2's pull architecture (§06.ii).
CASSCF	Complete Active Space Self-Consistent Field	A multireference quantum chemistry method that treats a chosen subset of orbitals (the active space) with full configuration-interaction accuracy. Cited as one of the classical baselines: CASSCF becomes intractable beyond approximately 18 active electrons, which is exactly why the four target mutations (TP53, KEAP1, CDKN2A, STK11) at 14-28 electrons sit at or beyond the classical ceiling.
CCSD(T)	Coupled Cluster with Single, Double, and perturbative Triple excitations	The classical "gold standard" of quantum chemistry for single-reference systems, with O(N ⁷) computational scaling. The poster lists CCSD(T) as infeasible at 44 active electrons and uses it as the reference value in the benchmarking comparison (~11.2 kcal/mol). VQE-SqDRIFT outputs are validated against published CCSD(T) values where available.
CDKN2A	Cyclin-Dependent Kinase Inhibitor 2A	A tumor-suppressor gene encoding p16 and p14ARF. CDKN2A loss is one of the four target mutations in the dissertation. Its hydrophobic core disruption requires an active space in the 14-18 electron range, straddling the

		classical tractability threshold and providing a less metallic counterpoint to TP53 in the dissertation's mutation portfolio.
CMC	Chemistry, Manufacturing, and Controls	The section of an FDA Investigational New Drug (IND) submission documenting drug composition, synthesis, and quality control. Out of dissertation scope (Part 312 territory) but mentioned in context to clarify the boundary between this dissertation's Part 11 contribution and the broader regulatory pathway.
CNN	Convolutional Neural Network	A deep-learning architecture specialized for grid-structured data such as images. Nvidia Ising's QEC decoder is implemented as a 3D CNN that acts as a fast pre-decoder across a three-dimensional syndrome lattice, resolving common errors in real time and escalating the remainder to a conventional global decoder. Referenced in Layer 2 of the architecture; the provenance of both stages is captured in field P9 of the audit trail (a forward-looking field — Phase 3 uses deterministic decoding).
CRO	Contract Research Organization	An external company providing outsourced research services to pharma. Referenced in the BPR analysis (problem P3) - current CRO handoffs proceed without formal go/no-go thresholds even when confidence intervals are low, illustrating the governance gap that the orchestration platform addresses.
CSA	Computer Software Assurance	FDA guidance (finalized September 2025) requiring risk-based validation of software used in regulated contexts. The dissertation's Layer 5 framework maps each pipeline component (Qiskit compiler, SqDRIFT sampler, AI scoring module, uncertainty quantifier) to CSA-compliant validation documentation.
CUDA-Q	Nvidia's quantum-classical programming model	A programming framework enabling unified development across GPU and QPU resources. Listed in Layer 2 (commodity substrate) of the architecture. Its emergence as a vendor-spanning standard is what reframed the dissertation's claim from "cross-vendor abstraction" (now commodity) to "Part-11-governed, bifurcation-aware orchestration" (defensible white space).
DBD	DNA-Binding Domain	The structural region of p53 (residues ~94–312) responsible for sequence-specific DNA recognition. C275 lies within the DBD; its mutation to Phe disrupts the loop-sheet-helix DNA-contact interface. The DBD is the spatial anchor for active-space selection in the dissertation.
DeepChem	DeepChem (open-source ML library for chemistry)	A Python toolkit providing pre-built deep-learning models for drug discovery, including Graph Neural Networks. Used in Layer 2 for AI rescoring of quantum simulation outputs - interpreting and ranking results by clinical relevance after the quantum kernel has produced raw binding-energy estimates.
DFT	Density Functional Theory	A workhorse classical method for electronic-structure calculations. Used in Layer 2 for preprocessing - determining the active space (which electrons enter the quantum simulation, which remain classical) before the quantum kernel runs. Also serves as a benchmark in the results comparison: at C275F's active-space boundary, DFT-B3LYP misses chemical accuracy by 2.8 kcal/mol.
DMRG	Density Matrix Renormalization Group	A classical numerical method exceptionally accurate for one-dimensional or weakly entangled systems. Listed in Layer 2 alongside DFT as a classical pre/post-processing tool: provides classical baselines and treats fragments outside the quantum active space.
DMRG-SCF	Density Matrix Renormalization Group Self-Consistent Field	A CASSCF calculation in which the DMRG sweep, rather than full configuration interaction, performs the orbital optimization. Because FCI's cost is combinatorial in the number of active orbitals and becomes impractical past roughly sixteen, DMRG-SCF is the step that lets a classifier realize DMRG's own premise of handling a large active space at polynomial rather than exponential cost (\$09).
DSR	Design Science Research	The methodological framework anchoring this dissertation, after Hevner et al. (2004). DSR positions the work as the design and evaluation of an artifact (the platform) rather than as hypothesis-testing science. Phase 4 implements DSR-prescribed expert evaluation with an n=4–6 panel, OSF pre-registered, against four criteria (utility, novelty, rigor, feasibility).
EGFR	Epidermal Growth Factor Receptor	An oncogene with FDA-approved targeted therapies for NSCLC (osimertinib, erlotinib). Referenced in the Problem Statement to delineate the clinical gap: ~60% of NSCLC patients lack EGFR/ALK/KRAS-G12C mutations and therefore have no approved targeted therapy - the patient population this dissertation addresses.
EWf-SQD	Embedded Wave Function - Sample-based Quantum Diagonalization	A hybrid quantum-classical workflow announced by IBM and Cleveland Clinic in March 2026, demonstrated on the 303-atom Trp-cage mini-protein. Combines wave-function-based fragment embedding with SQD on each cluster. Recognized late in the dissertation cycle as an adjacent methodology with potential application to scaffold-perturbation mutations like C275F; its convergence guarantees are weaker than SqDRIFT's, which becomes a design consideration if the methods are hybridized.
FastAPI	Python asynchronous web framework	The framework implementing the SOLANGE backend — dispatch queue, LEON notarization endpoints, and the Computational Decision Gateway. Deployed on Render and named in the deployment topology (\$06.ii).
FCI	Full Configuration Interaction	The exact solution to the electronic Schrödinger equation within a given basis set. Computationally infeasible beyond very small systems: $O(2^N)$ scaling reaches $\sim 10^{13}$ operations at $N=44$, which the poster cites as completely intractable. FCI defines the unreachable upper bound that quantum methods aim to approximate.
FDA	U.S. Food and Drug Administration	The federal agency regulating pharmaceuticals, medical devices, and clinical research in the United States. This dissertation's Layer 5 contribution operates within FDA's 21 CFR Part 11 framework (electronic records and signatures) - designing audit-trail architecture for probabilistic quantum simulation outputs.
FEP+	Free Energy Perturbation Plus (Schrödinger's commercial implementation)	A leading classical method for computing protein-ligand binding affinities. Cited in Layer 3 (benchmarking) as the industrial classical baseline against which quantum methods must be compared in terms of wall-clock time, accuracy, and active-space coverage.
GARD	Genomic-Adjusted Radiation Dose	A radiogenomic model predicting individual tumor radiosensitivity from gene-expression signatures. Mentioned in the Future Research Direction on the poster: post-doctoral work would extend the dissertation's quantum-derived mechanistic features into GARD-class radiotherapy decision support for TP53-mutant tumors. Out of dissertation scope.
GNN	Graph Neural Network	A neural-network architecture operating on graph-structured data - natural for molecules represented as atoms (nodes) and bonds (edges). DeepChem GNN is the AI rescoring component in Layer 2, interpreting raw quantum-simulation outputs in terms of drug-likeness and clinical relevance.

GPU	Graphics Processing Unit	A massively parallel processor originally designed for graphics, now central to scientific computing. NVQLink provides the GPU ↔ QPU interconnect at the heart of Layer 2; classical pre/post-processing (DFT, DMRG, GNN inference) runs on GPU clusters.
Hartree-Fock	HF — mean-field electronic-structure method	The single-determinant approximation in which each electron experiences only the averaged field of the others. It is the reference state on which the executed quantum-hardware runs measure $\langle H \rangle$ (§07.4), and the zero-entanglement baseline against which every correlated method is defined.
HIPAA	Health Insurance Portability and Accountability Act	U.S. federal law (1996) governing privacy of medical information. The dissertation works exclusively with de-identified single-case data (MI25-0349) to remain outside HIPAA's PHI requirements. HIPAA compliance for multi-patient analysis is reserved for post-doctoral work.
HPC	High-Performance Computing	Large-scale classical computing infrastructure (CPU/GPU clusters). The third pillar of the QC·AI·HPC architecture: HPC handles classical pre-processing (DFT, DMRG) and post-processing (AI rescoring) that surrounds the quantum kernel.
IND	Investigational New Drug	An FDA application (governed by 21 CFR Part 312) required before administering an experimental drug to humans. Out of dissertation scope by explicit decision: Part 312 is a regulatory pathway, while this dissertation's contribution is the upstream Part 11 audit infrastructure that any future IND submission relying on quantum-simulation evidence would require.
IRB	Institutional Review Board	The ethics committee approving research involving human subjects. The dissertation operates under an IRB scope limited to anonymized single-case reference (MI25-0349); multi-patient analysis is deferred to post-doctoral work.
IS	Information Systems	The disciplinary home of this dissertation. Positioning the work as IS - architecture, governance, audit-trail design - rather than as computational chemistry is what makes it appropriate for CGU CISAT's DTech program.
JW	Jordan-Wigner (fermion-to-qubit mapping)	One of three fermion-to-qubit mappings the Phase 3A pipeline implements and compares (alongside parity+two-qubit-reduction and Bravyi-Kitaev). The mapper choice is recorded in audit-trail field P2 (compilation lineage).
KEAP1	Kelch-like ECH-Associated Protein 1	One of the three KEAP1 target mutations (LOF, G333C, R320Q), classified Class B (quantum-advantaged). The Kelch-domain binding pocket — which mediates the Nrf2 interaction disrupted in cancer — spans a full active site of ~155 active electrons (~310 qubits under Jordan-Wigner encoding), well beyond the exact-classical wall (~18e). Designated fault-tolerant QPU target (~2030+); coordinate-verified from PDB structures 1U6D and 2FLU.
Kelch	Kelch domain — KEAP1's substrate-binding β-propeller	The KEAP1 domain mediating the Nrf2 interaction disrupted in cancer. Its binding pocket defines KEAP1's ~155-electron active site, and the fact that loss of function arises across the pocket rather than at one residue is why KEAP1 carries no single-residue anchor and is classified at gene level (§06.i).
KRAS	Kirsten Rat Sarcoma viral oncogene homolog	An oncogene mutated in ~25% of NSCLC. KRAS G12C now has an FDA-approved targeted therapy (sotorasib), placing it among the druggable mutations referenced in the Problem Statement. Most other KRAS variants and the four mutations addressed in this dissertation remain non-druggable.
L2 / L3 (loops)	L2 and L3 loops of the p53 DNA-binding domain	Structural loops in the p53 DBD that stabilize Zn ²⁺ tetrahedral coordination via Cys176, His179, Cys238, Cys242. The loop-sheet-helix motif containing C275 is structurally adjacent. Relevant to active-space selection: the local 24-electron Phase 3A space and the full 44-electron Phase 3B space differ in how much of L2/L3 they include.
Laguna	USC CARC compute cluster	The SLURM-managed cluster on which the classical HPC and DMRG rungs of the orchestration ladder execute. Its prohibition on inbound connections is the security boundary the bounded execution agent runs inside, and therefore the environment DP2's pull architecture is designed against (§06.ii).
ML	Machine Learning	A subfield of AI focused on systems that learn patterns from data. The Nvidia Ising 3D CNN QEC decoder is an ML component; its provenance - model version, weights, calibration data - is captured in audit-trail field P9 (ML-decoder lineage).
MM	Molecular Mechanics	Classical force-field treatment of atomic motion. The classical half of QM/MM hybrid partitioning; the bulk of the protein is treated by MM while the chemically critical region around the mutation site receives quantum treatment.
MTD	Maximum Tolerated Dose	The highest drug dose that does not produce unacceptable toxicity. Mentioned in clinical-trial context (e.g., the rezatapopt PYNACLE Phase 1 endpoint). Out of dissertation scope; relevant only to clinical-translation discussions in the Future Research section.
NDA	New Drug Application	The FDA filing seeking approval to market a new drug. Out of dissertation scope; mentioned in connection with rezatapopt (PMV Pharma's NDA targeted for Q1 2027 for TP53 Y220C) to clarify why Y220C is no longer non-druggable.
NEJM	New England Journal of Medicine	A leading medical journal. Cited as the venue (February 2026) for the rezatapopt Phase 1 results in TP53 Y220C - the publication that confirms Y220C's transition from non-druggable to druggable, prompting its removal from the dissertation's mutation list.
Netlify	Static-site hosting and continuous-deployment service	Hosts the SOLANGE front end and the deployed dissertation, deploying automatically from the repository's main branch. Appears in the deployment topology alongside Render and Supabase (§06.ii).
NGS	Next-Generation Sequencing	High-throughput DNA sequencing technology used to identify mutations in patient tumors. The dissertation's anchor patient case MI25-0349 is sourced from NGS data: TP53 p.Cys275Phe identified as a Tier I pathogenic variant at 12.66% allele frequency (Emek Medical Center, June 2025).
NSCLC	Non-Small Cell Lung Cancer	The most common form of lung cancer (~85% of cases) and the leading cause of cancer mortality globally (~2.1 million new cases and ~1.53 million deaths annually; GLOBOCAN 2022). The clinical context of the dissertation: the four target mutations (TP53, KEAP1, CDKN2A, STK11) are common non-druggable drivers in NSCLC patients.
LEON	Lineage-Evidence Orchestration & Notarization	The notarization component on the return path of the pull-based dispatch loop (§06.ii). LEON independently recomputes the P8 cryptographic seal and re-checks physics-consistency at ingestion, rejecting any externally-submitted run whose lineage does not verify, and re-attests integrity on demand at query time. It operationalizes the verify, don't trust principle — trust derived from reproducible cryptographic proof, not from the identity of the boundary a result arrived through — and is the element of the architecture with no counterpart in the precedent systems of §3.4. Named in memory of the candidate's father, Leon.

Nvidia	Nvidia Corporation	GPU manufacturer providing key Layer 2 components: CUDA-Q programming model, NVQLink GPU ↔ QPU interconnect, and Ising AI stack (VLM calibration + 3D CNN QEC decoder). Nvidia's emergence as cross-vendor infrastructure is what made the dissertation's earlier "cross-vendor abstraction" claim a commodity, prompting the narrowing to Part-11 governance as the defensible white space.
NVQLink	Nvidia's GPU ↔ QPU interconnect	A hardware-software fabric enabling tight integration between GPU clusters and quantum processors. Layer 2 component (commodity); together with CUDA-Q, it represents the vendor-spanning quantum-classical infrastructure the dissertation builds on top of rather than reinventing.
OSF	Open Science Framework	A platform for pre-registration of research designs and hypotheses. The dissertation's Phase 4 evaluation criteria are pre-registered on OSF before evaluation begins - the methodological commitment that ensures Layer 1 acceptance thresholds cannot be retrofitted to results.
P1- P9	Provenance fields 1 through 9 of the audit-trail schema	The dissertation's signature contribution to Information Systems: a nine-field provenance record that makes probabilistic quantum simulation outputs auditable under FDA 21 CFR Part 11. P1 circuit fingerprint, P2 compilation lineage, P3 device identifier and calibration epoch, P4 per-gate error budget snapshot, P5 shot count and raw outcome distribution, P6 error-mitigation lineage, P7 statistical estimator and confidence interval, P8 cryptographic seal, P9 ML-decoder provenance. Each field maps to a specific 21 CFR §11.10 requirement.
Part 11	21 CFR Part 11 (FDA Electronic Records & Signatures rule)	The federal regulation governing electronic records and signatures in FDA-regulated contexts. The dissertation's Layer 5 / §06.iii contribution: the first systematic framework applying Part 11 to probabilistic quantum-computed binding data. Part 11 (records) is in scope; Part 312 (IND submission) is explicitly out of scope.
Part 312	21 CFR Part 312 (FDA Investigational New Drug regulation)	The federal regulation governing IND applications. Explicitly excluded from the dissertation's novelty claim - Part 312 is a regulatory pathway, while the contribution sits upstream in the Part 11 records infrastructure that any future IND filing using quantum evidence would depend on.
PennyLane	PennyLane (Xanadu's quantum-machine-learning framework)	The simulator backend for Phase 3A. Currently live at 4 qubits (CAS(2e,2o)/STO-3G, AllSinglesDoubles UCCSD ansatz, 80 Adam steps) on Render cloud. The same PennyLane pipeline scales to the intermediate target (~20e/40q, ~17.6 TB state-vector) on NSF ACCESS / USC CARC HPC, enabling architecture validation at meaningful quantum-circuit scale without IBM Quantum Network access before Phase 3B hardware execution. The scientific ceiling target is ~24e/48q (~4.5 PB state-vector, edge of global classical capacity); the operational run uses ~20e/40q, which is well within multi-node HPC reach while still demonstrating two orders of magnitude of scale beyond the 4-qubit minimal prototype.
PHI	Protected Health Information	Patient-identifying medical data protected under HIPAA's eighteen-identifier rule. The dissertation works only with de-identified data (MI25-0349 is an internal study identifier with no link to patient identity), placing it outside PHI/HIPAA scope and simplifying compliance dramatically.
PMV	PMV Pharmaceuticals	The biotech company developing rezatapopt (PC14586) for TP53 Y220C. PMV's Phase 2 PYNNAACLE trial and planned Q1 2027 NDA are the reason Y220C transitioned from non-druggable to druggable and was removed from the dissertation's mutation list.
PYNNAACLE	Phase 1/2 trial of rezatapopt in TP53 Y220C-mutated solid tumors (NCT04585750)	The clinical trial demonstrating rezatapopt's efficacy in Y220C-mutated cancers (33% ORR across 8 tumor types as of August 2025). Cited as the evidence base for removing Y220C from the dissertation's non-druggable-mutation list.
PySCF	Python-based Simulations of Chemistry Framework	The open-source quantum-chemistry library behind the platform's classical electronic-structure work: mean-field solutions, active-space construction, and the CASSCF/FCI reference energies against which computed results are checked.
QC	Quantum Computing	Computation exploiting quantum-mechanical phenomena (superposition, entanglement) to solve problems intractable for classical hardware. The first pillar of the QC-AI-HPC architecture; Layer 1 of the platform houses the quantum kernel (SqDRIFT on IBM Heron r3).
QCS	Quantum Cloud Services (IBM)	IBM's cloud platform for quantum-hardware access. Listed alongside Nvidia NVQLink in the bifurcation-aware orchestration claim: the platform's architecture must survive the IBM-QCS vs Nvidia-NVQLink vendor split that emerged during the proposal cycle.
QEC	Quantum Error Correction	The set of techniques for protecting quantum computations against noise by encoding logical qubits across many physical qubits with active error detection. The Nvidia Ising 3D CNN decoder is the QEC decoder in Layer 2; its operation generates the data captured in audit-trail field P9 (ML-decoder lineage).
QEM	Quantum Error Mitigation	Techniques (ZNE, probabilistic error cancellation, readout correction) that reduce - but do not eliminate - the impact of noise on quantum computations, without the qubit overhead of full QEC. Recorded in audit-trail field P6 (error-mitigation lineage). Distinguished from QEC: mitigation is post-processing; correction operates during execution.
Qiskit	IBM's open-source quantum computing SDK	The stack through which quantum circuits are constructed, transpiled, and submitted to IBM hardware. Qiskit Runtime supplies the production primitives for fermionic Hamiltonian construction, sample-state preparation, and classical diagonalization on which the Phase 3B SqDRIFT path depends (§02).
QM	Quantum Mechanics	The physics framework underlying quantum chemistry and quantum computing. The quantum half of QM/MM hybrid partitioning: the chemically critical region around the mutation site receives QM treatment while the bulk of the protein is treated classically by MM.
QM/MM	Quantum Mechanics / Molecular Mechanics hybrid partitioning	A multi-scale modeling approach in which a small chemically critical region is treated quantum-mechanically while the surrounding environment is treated classically. The dissertation requires a novel QM/MM scheme tailored to TP53's Zn ²⁺ site for C275F - itself a research contribution within the broader pipeline.
QPU	Quantum Processing Unit	The hardware unit performing quantum computation, analogous to a CPU or GPU. IBM Heron r3 is the QPU targeted in Phase 3B. NVQLink provides the GPU ↔ QPU interconnect.
QV	Quantum Volume	An IBM-defined benchmark capturing combined qubit count, connectivity, and gate fidelity. The poster lists IBM Heron r3 at QV ≥ 2048 - establishing that the hardware regime of the IBM Research 2026 anchor (32 electrons, 100 qubits) is reproducible at Phase 3B scale.
RCSB PDB	Research Collaboratory for Structural	The public repository of experimentally determined macromolecular structures, and the source of the coordinate files behind the target registry (e.g. 2OCJ, the p53 DNA-binding domain at 2.05 Å). Also the basis of the observation that no

	Bioinformatics Protein Data Bank	crystallographic structure of the C275F mutant itself has been deposited (§03).
RP2D	Recommended Phase 2 Dose	The drug dose level chosen from Phase 1 results to take into Phase 2 trials. Mentioned in clinical-trial context (rezatapopt PYNACLE) and out of dissertation scope; relevant only to the rationale for removing Y220C from the mutation list.
sbatch	SLURM batch-submission command	The command by which a researcher submits a non-interactive job to the cluster queue. It is one of the two entry modes the DMRG classification rung admits — the other being an execution agent claiming work from the platform's dispatch queue (§06.ii, Figure 6.2).
SC150	Stabilization Concentration 150 (50% protein-stabilization endpoint)	A biochemical potency measure for protein-stabilizing drugs. Cited in the rezatapopt context (SC150 = 9 nM for Y220C binding) - the kind of experimental ground-truth value that, were Y220C retained as a positive control, the dissertation's quantum platform could be benchmarked against.
SLURM	Simple Linux Utility for Resource Management	The batch scheduler allocating compute nodes on the Laguna cluster. Its queue is the boundary at which the orchestration layer hands a classical job over to institutional infrastructure and waits for a sealed provenance record to come back (§08).
SQD	Sample-based Quantum Diagonalization	A family of hybrid quantum-classical algorithms where the quantum computer samples configurations and the classical computer diagonalizes a reduced Hamiltonian over those samples. SqDRIFT is one variant; EWF-SQD is another. Convergence in the SQD family is generally empirical rather than provable.
SqDRIFT	Sample-based quantum Diagonalization with stochastic Trotter (qDRIFT)	The quantum algorithm anchoring Layer 1 of the dissertation. Demonstrated by IBM in 2026 (IBM Research; arXiv:2603.08696 and the SqDRIFT method paper arXiv:2508.02578): a 32-electron active space of $C_{13}Cl_2$ on 100 qubits of Heron, beyond the classical CASSCF ceiling of ~18 electrons. SqDRIFT inherits qDRIFT's provable convergence bounds - a strength the dissertation contrasts with EWF-SQD's empirical-only convergence.
STK11	Serine/Threonine Kinase 11 (also LKB1)	One of the four STK11/LKB1 target mutations (LKB1 LOF, F354L, D194N), classified Class A (quantum-necessary). Full active site: ~152 active electrons (~304 qubits under Jordan-Wigner encoding) — designated fault-tolerant QPU target (~2030+). STK11 F354L local site (~48e/~96q; AlphaFold Q15831) is a near-term exception executable on IBM Heron r3, only 2 qubits beyond the demonstrated 94-qubit ceiling (Merz et al. 2026, arXiv:2605.01138). PDB source: 2WTK.
TP53	Tumor Protein 53 (the p53 tumor suppressor gene)	The most frequently mutated gene in human cancers (~59% of solid tumors). The dissertation focuses on TP53 structural (DNA-interface) mutants - particularly C275F - that disrupt DNA-binding geometry without creating druggable pockets. TP53's Zn^{2+} tetrahedral coordination site (Cys176/His179/Cys238/Cys242) requires ~22–28 active electrons, the paradigmatic case for SqDRIFT-grade simulation.
Supabase	Hosted PostgreSQL platform with authentication and row-level security	The database and identity layer behind the platform: P1–P9 provenance records, the LEON audit table, the dispatch queue, and per-user authentication. Its append-only audit policy — insert and read permitted, update and delete denied — is what makes the Part 11 claim testable in running code rather than asserted (§06.iii).
Trp-cage	Tryptophan-cage miniprotein	A 20-residue, 303-atom synthetic miniprotein. The substrate of the Cleveland Clinic / IBM EWF-SQD demonstration (March 2026) - the first protein-scale SQD-family simulation. Used in the dissertation's positioning to argue that EWF-SQD validates protein-scale viability without scooping the dissertation's clinical-mutation focus.
VLM	Vision-Language Model	A large multimodal AI model processing images and text jointly. Nvidia's Ising calibration component (35-billion-parameter VLM) autonomously brings up QPU calibration. Layer 2 commodity; its outputs feed the QEC decoder whose lineage is captured in P9.
VQE	Variational Quantum Eigensolver	A hybrid quantum-classical algorithm that uses a classical optimizer to minimize the expectation value of a quantum-prepared trial state. Phase 3A benchmarks VQE outputs against DFT and CCSD(T) reference values on the C275F local active site. The poster's results illustration shows VQE-SqDRIFT at -11.7 kcal/mol with 0.5 kcal/mol error - within chemical accuracy.
Y220C	TP53 p.Tyr220Cys (tyrosine-to-cysteine substitution at position 220)	A TP53 missense mutation creating a small drug-binding pocket. Originally listed among the dissertation's five target mutations but removed once rezatapopt (PMV Pharma's first-in-class p53 reactivator) entered Phase 2 with FDA Fast Track designation, reclassifying Y220C as druggable. The current target list is four mutations: C275F, KEAP1, CDKN2A, STK11.
ZNE	Zero-Noise Extrapolation	A quantum error-mitigation technique that runs circuits at varying noise levels and extrapolates to the zero-noise limit. Listed among the QEM techniques whose vendor-specific implementation differences (BPR problem P2) the orchestration platform must standardize. Recorded in audit-trail field P6.

Note on scope. This glossary is descriptive of the work as it stands at the proposal stage. Some terms (e.g., EWF-SQD, ZNE) appear in the dissertation only as comparison points or discussion items, while others (e.g., SqDRIFT, BQP, Part 11, P1–P9) are core to the contribution. The "Definition & Context" column makes this distinction explicit for each entry.

Living document. As the dissertation progresses through Phases 3A and 3B, additional terms (specific calibration data fields, hardware identifiers, evaluation metrics) will enter the artifact set. This glossary will be updated alongside each major revision of the proposal or platform.